[301] Randomness

Tyler Caraza-Harter

Which series was randomly generated? Which did I pick by hand?

1

2

Announcement 1: Recommended popular stats books (for summer reading)

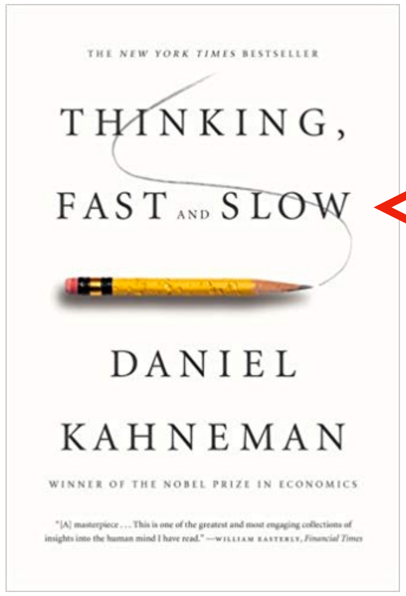

Thinking, Fast and S
by Daniel Kahnema

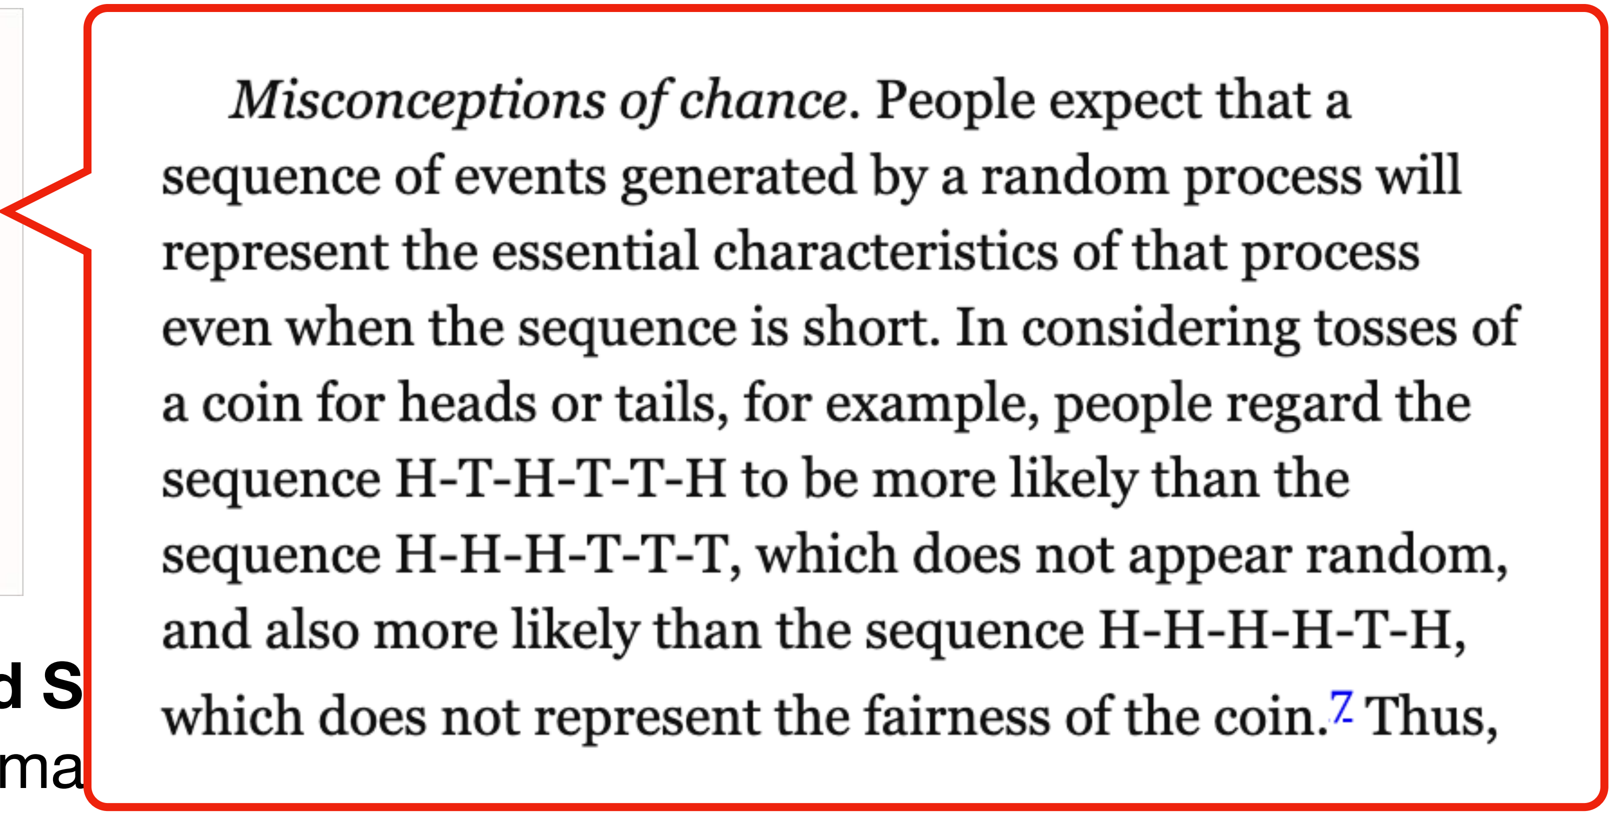

Misconceptions of chance. People expect that a sequence of events generated by a random process will represent the essential characteristics of that process even when the sequence is short. In considering tosses of a coin for heads or tails, for example, people regard the sequence H-T-H-T-T-H to be more likely than the sequence H-H-H-T-T-T, which does not appear random, and also more likely than the sequence H-H-H-H-T-H, which does not represent the fairness of the coin.[7] Thus,

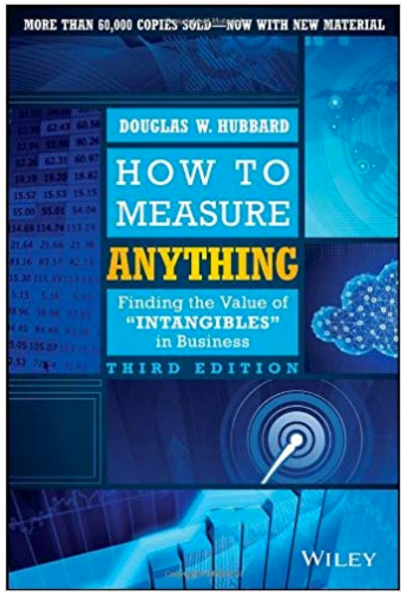

How to Measure Anything
by Douglas W. Hubbard

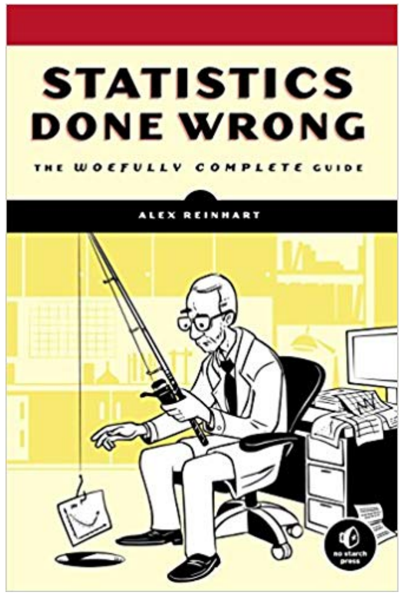

Statistics Done Wrong
by Alex Reinhart

Announcement 1: Recommended popular stats books (for summer reading)

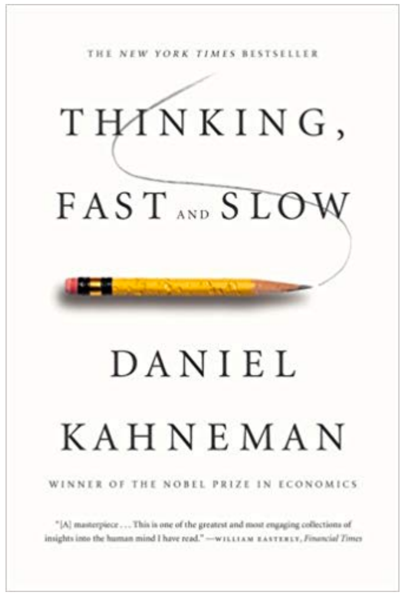

Thinking, Fast and Slow
by Daniel Kahneman

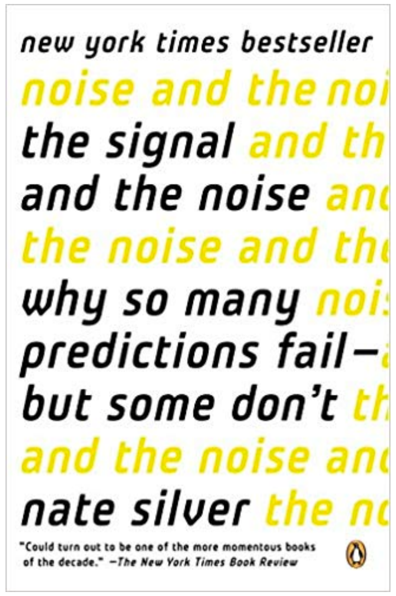

The Signal and the Noise
by Nate Silver

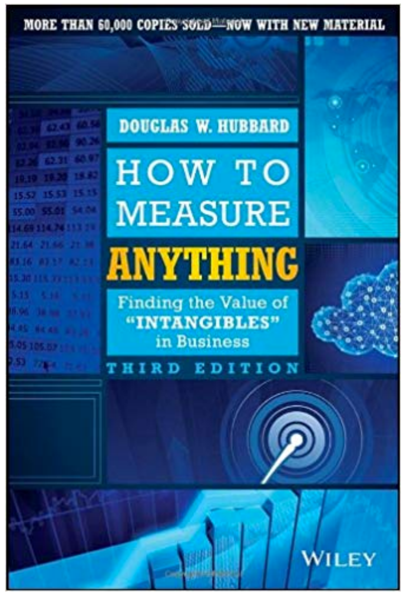

How to Measure Anything
by Douglas W. Hubbard

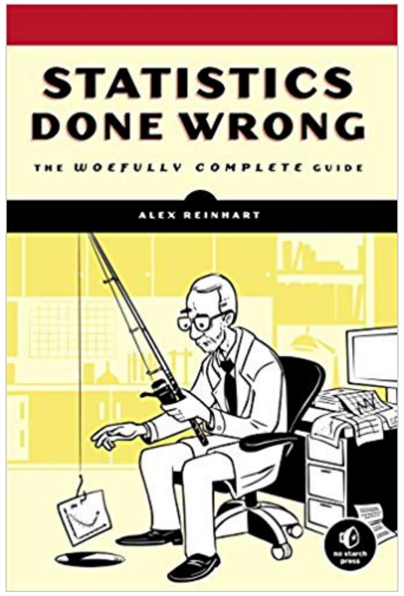

Statistics Done Wrong
by Alex Reinhart

Announcement 2: Course Evaluations

Section 2:

https://aefis.wisc.edu/index.cfm/page/AefisCourseSection.surveyResults?courseSectionid=593535

Section 3:

https://aefis.wisc.edu/index.cfm/page/AefisCourseSection.surveyResults?courseSectionid=593536

I always read all the feedback, so please take the time to complete these!

Announcement 3: Final Exam Prep

Details: similar to midterms

- worth 20%
- 2 hours on May 8th at 7:45am (in the morning!)
- you can have a single page of notes (both sides), as usual
- cumulative, across whole semester
- prep for Friday review session
- **watch your email for room details!**

Recommended prep

- make sure you understand all the worksheet problems
- review the readings, especially anything I took the time to write myself
- review everything you got wrong on the midterms
- review the slides
- review the code you wrote for the projects

Things not on the old final that we covered this semester

- beautifulsoup
- randomness

Why Randomize?

Games

Security

Simulation

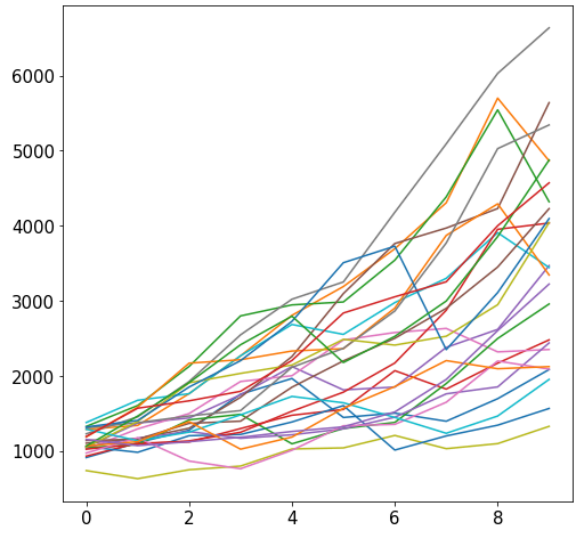

Why Randomize?

Games

Security

Simulation

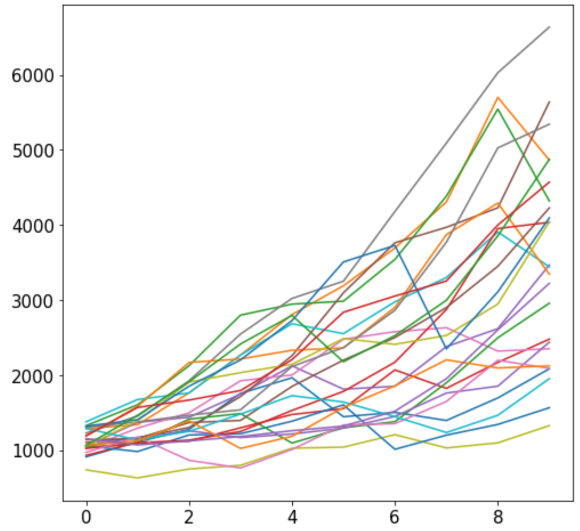

our focus

Outline

choice()

bugs and seeding

significance

histograms

normal()

New Functions Today

`numpy.random`:

- powerful collection of functions
- **`choice`**, **`normal`**

`Series.line.hist`:

- similar to bar plot
- visualize spread of random results

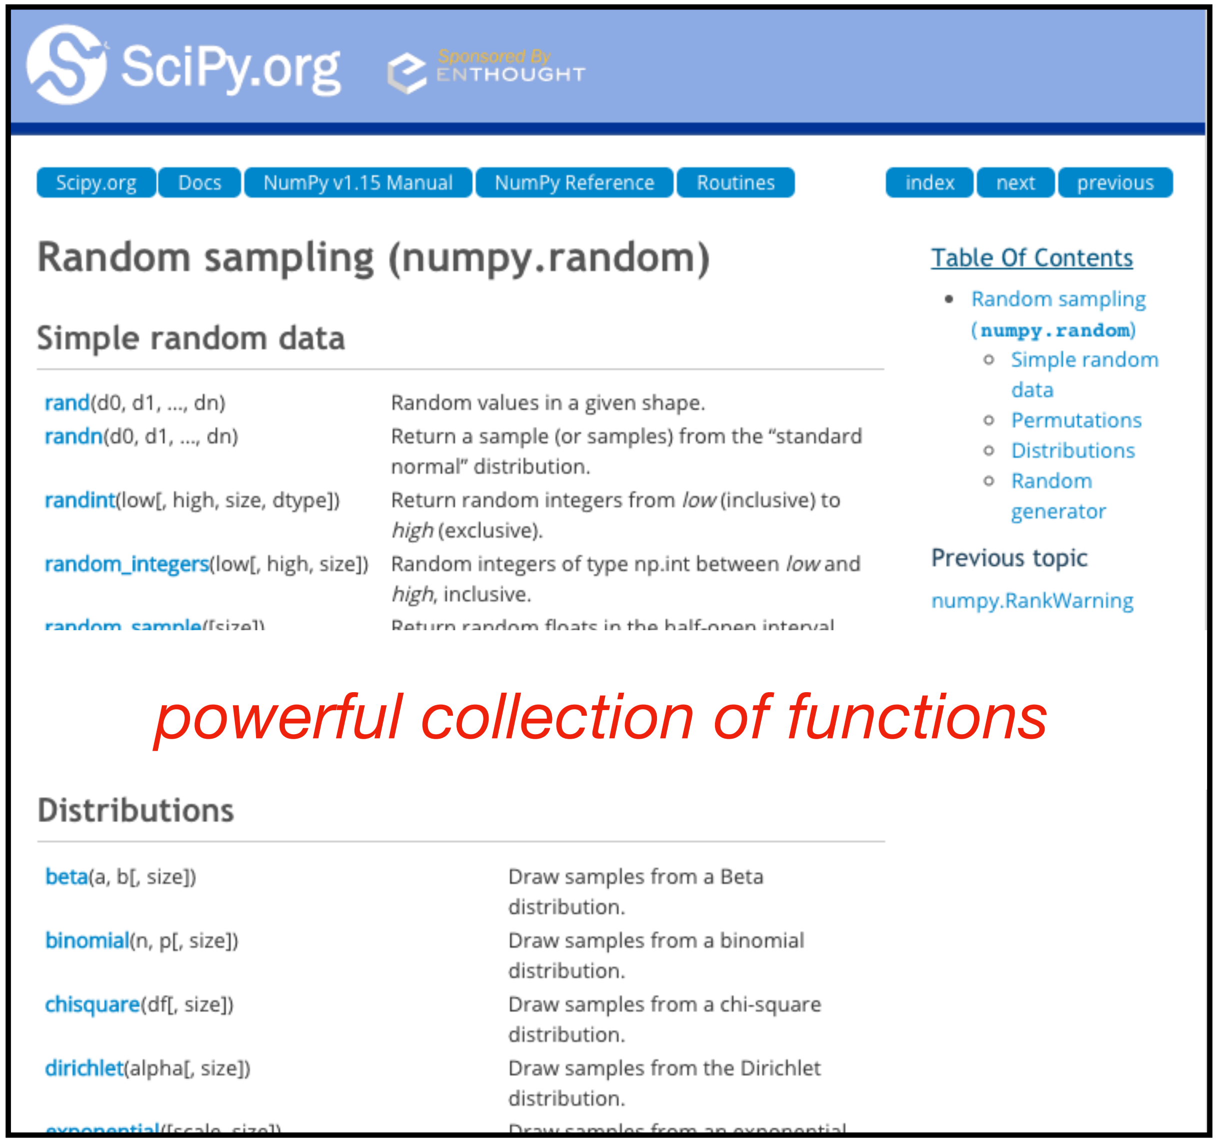

Random sampling (numpy.random)

Simple random data

rand(d0, d1, ..., dn)	Random values in a given shape.
randn(d0, d1, ..., dn)	Return a sample (or samples) from the "standard normal" distribution.
randint(low[, high, size, dtype])	Return random integers from *low* (inclusive) to *high* (exclusive).
random_integers(low[, high, size])	Random integers of type np.int between *low* and *high*, inclusive.

powerful collection of functions

Distributions

beta(a, b[, size])	Draw samples from a Beta distribution.
binomial(n, p[, size])	Draw samples from a binomial distribution.
chisquare(df[, size])	Draw samples from a chi-square distribution.
dirichlet(alpha[, size])	Draw samples from the Dirichlet distribution.

Table Of Contents

- Random sampling (`numpy.random`)
  - Simple random data
  - Permutations
  - Distributions
  - Random generator

Previous topic

numpy.RankWarning

choice

```
from numpy.random import choice, normal
```

choice

```
from numpy.random import choice, normal

result = choice(        )
```

list of things to randomly choose from

choice

```
from numpy.random import choice, normal

result = choice(["rock", "paper", "scissors"])
```

list of things to randomly choose from

https://www.securifera.com/blog/2015/09/09/mmactf-2015-rock-paper-scissors-rps/

choice

```
from numpy.random import choice, normal

result = choice(["rock", "paper", "scissors"])
print(result)
```

https://www.securifera.com/blog/2015/09/09/mmactf-2015-rock-paper-scissors-rps/

Output:

```
scissors
```

choice

```
from numpy.random import choice, normal

result = choice(["rock", "paper", "scissors"])
print(result)

result = choice(["rock", "paper", "scissors"])
print(result)
```

Output:

```
scissors
rock
```

choice

```
from numpy.random import choice, normal

result = choice(["rock", "paper", "scissors"])
print(result)

result = choice(["rock", "paper", "scissors"])
print(result)
```

Output:

```
scissors
rock
```

each time choice is called, a value is randomly selected (will vary run to run)

choice

```
from numpy.random import choice, normal

choice(["rock", "paper", "scissors"], size=5)
```

for simulation, we'll often want to compute many random results

choice

```
from numpy.random import choice, normal

choice(["rock", "paper", "scissors"], size=5)
```

```
array(['rock', 'scissors', 'paper', 'rock', 'paper'], dtype='<U8')
```

it's list-like

Random values and Pandas

```
from numpy.random import choice, normal

# random Series
Series(choice(["rock", "paper", "scissors"], size=5))
```

Random values and Pandas

```
from numpy.random import choice, normal

# random Series
Series(choice(["rock", "paper", "scissors"], size=5))
```

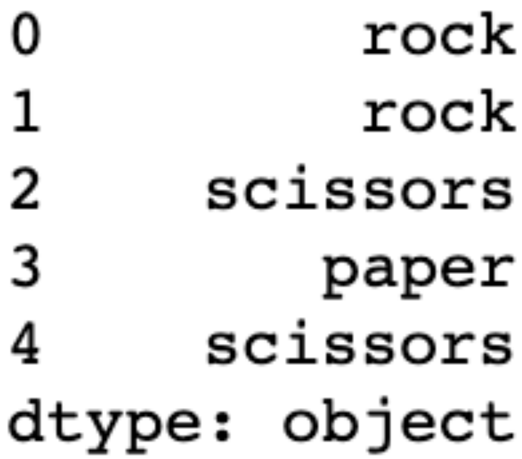

```
0        rock
1        rock
2    scissors
3       paper
4    scissors
dtype: object
```

Random values and Pandas

```
from numpy.random import choice, normal

# random Series
DataFrame(choice(["rock", "paper", "scissors"],
                 size=(5,2)))
```

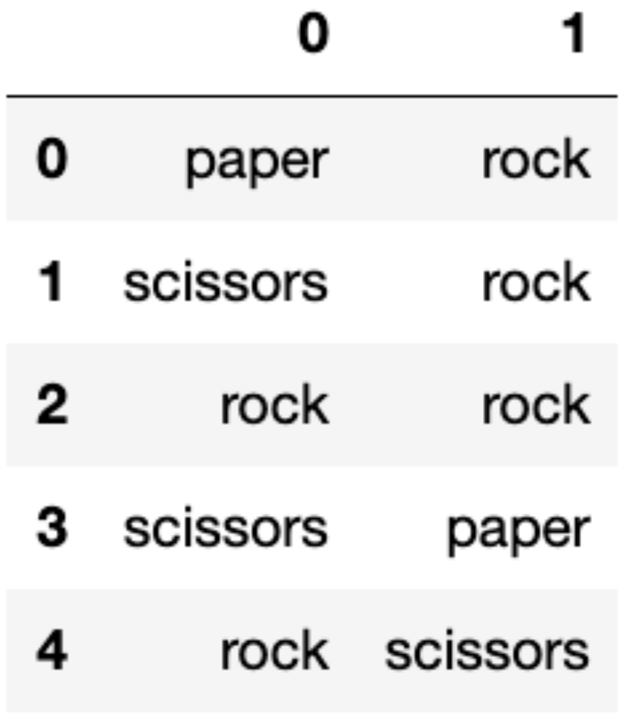

	0	1
0	paper	rock
1	scissors	rock
2	rock	rock
3	scissors	paper
4	rock	scissors

Demo: exploring bias

```
choice(["rock", "paper", "scissors"])
```

Question 1: how can we make sure the randomization isn't biased?

Demo: exploring bias

```
choice(["rock", "paper", "scissors"])
```

Question 1: how can we make sure the randomization isn't biased?

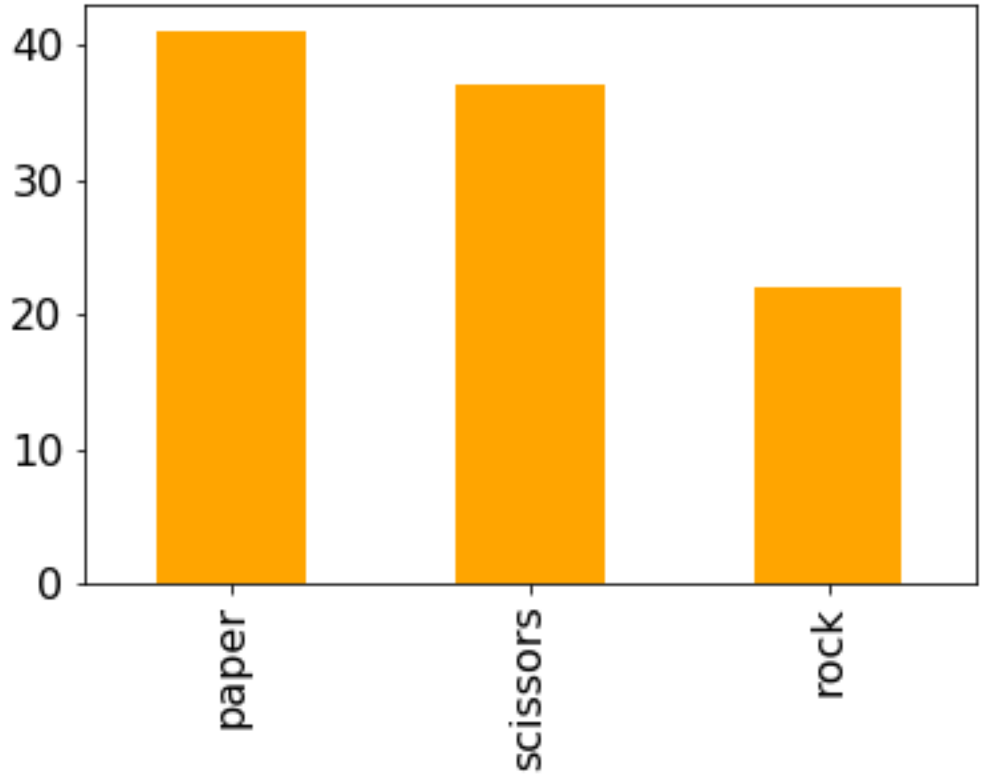

Demo: exploring bias

```
choice(["rock", "paper", "scissors"])
```

Question 1: how can we make sure the randomization isn't biased?

Question 2: how can we make it biased (if we want it to be)?

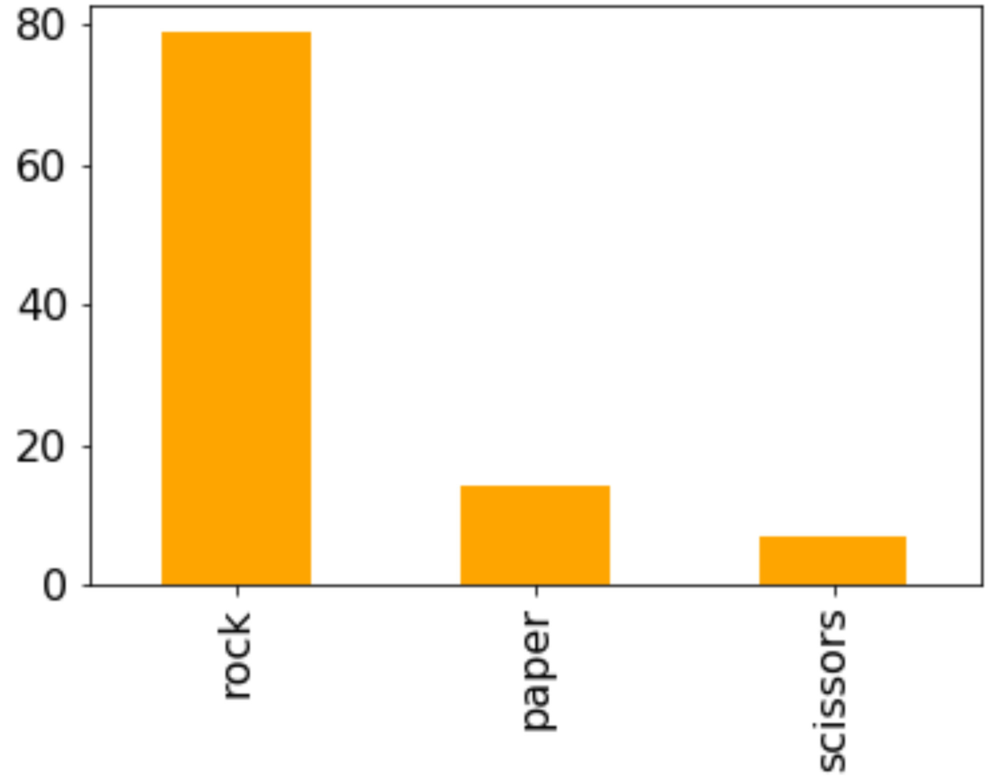

p=[...]

Random Strings vs. Random Ints

```
from numpy.random import choice, normal

# random string: rock, paper, or scissors
choice(["rock", "paper", "scissors"])
```

Random Strings vs. Random Ints

```
from numpy.random import choice, normal

# random string: rock, paper, or scissors
choice(["rock", "paper", "scissors"])

# random int: 0, 1, or 2
choice([0, 1, 2])
```

Random Strings vs. Random Ints

```
from numpy.random import choice, normal

# random string: rock, paper, or scissors
choice(["rock", "paper", "scissors"])

# random int: 0, 1, or 2
choice([0, 1, 2])
```

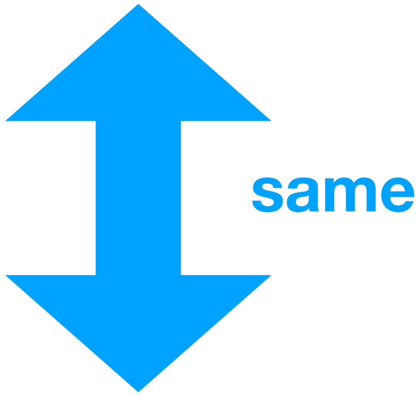

```
# random int (approach 2): 0, 1, or 2
choice(3)
```

Random Strings vs. Random Ints

```
from numpy.random import choice, normal

# random string: rock, paper, or scissors
choice(["rock", "paper", "scissors"])

# random int: 0, 1, or 2
choice([0, 1, 2])
```

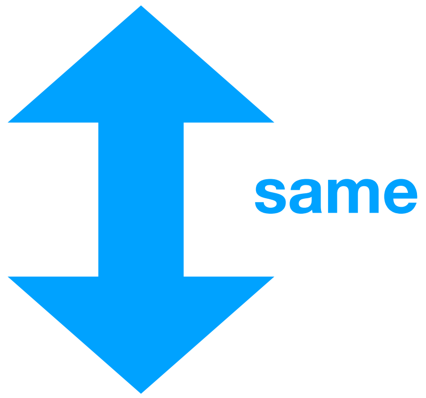

```
# random int (approach 2): 0, 1, or 2
choice(3)
```

random non-negative int
that is **less than 3**

Outline

choice()

bugs and seeding

significance

histograms

normal()

Example: change over time

```
s = Series(choice(10, size=5))
```

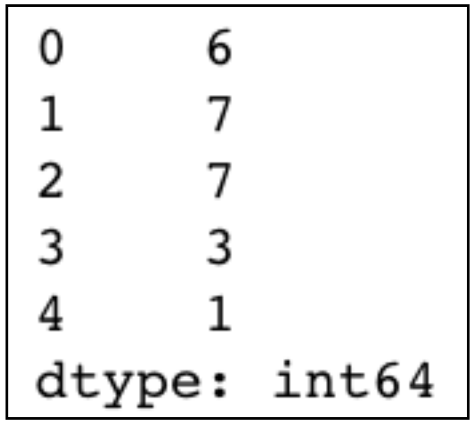

```
0    6
1    7
2    7
3    3
4    1
dtype: int64
```

```
s.plot.line()
```

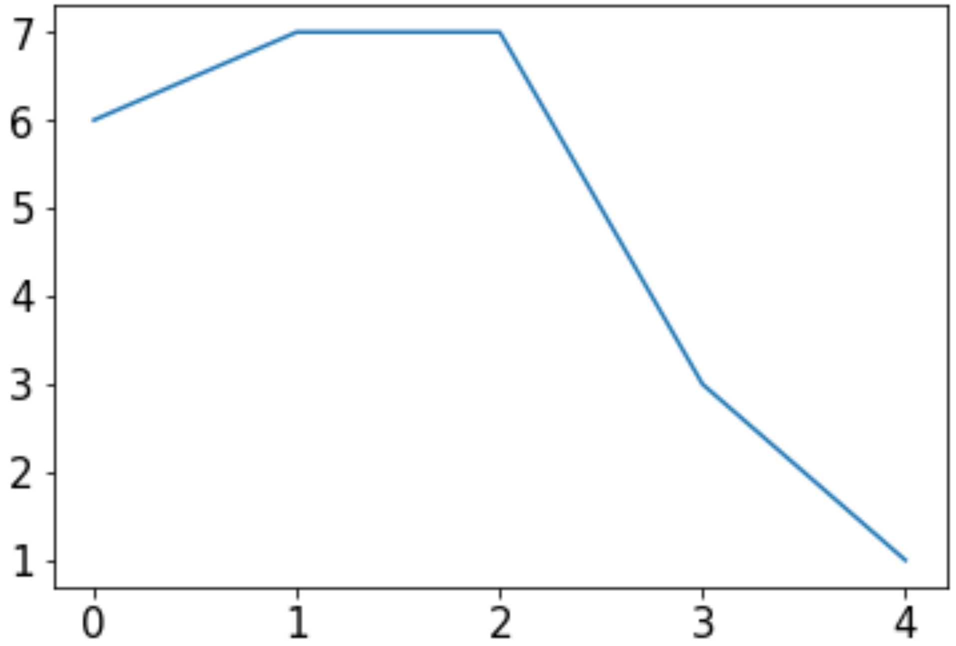

Example: change over time

```
s = Series(choice(10, size=5))
```

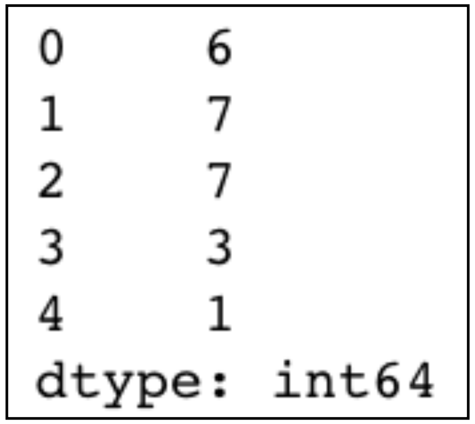

```
0     6
1     7
2     7
3     3
4     1
dtype: int64
```

```
s.plot.line()
```

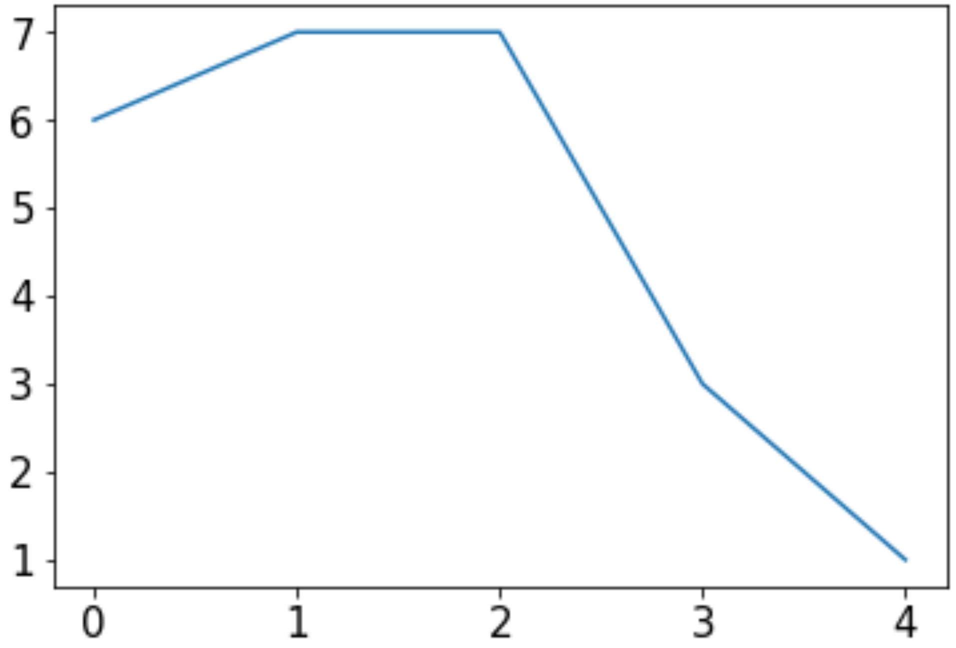

```
percents = []
for i in range(1, len(s)):
    diff = 100 * (s[i] / s[i-1] - 1)
    percents.append(diff)
Series(percents).plot.line()
```

what are we computing for diff?

Example: change over time

```
s = Series(choice(10, size=5))
```

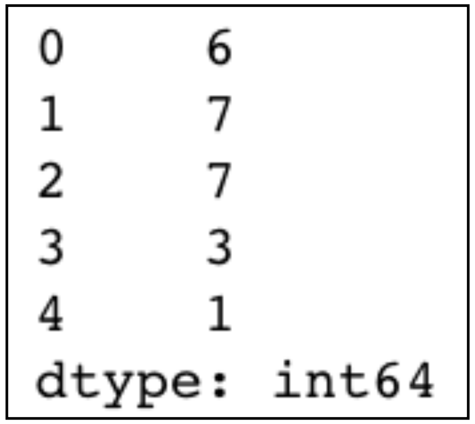

```
0      6
1      7
2      7
3      3
4      1
dtype: int64
```

```
s.plot.line()
```

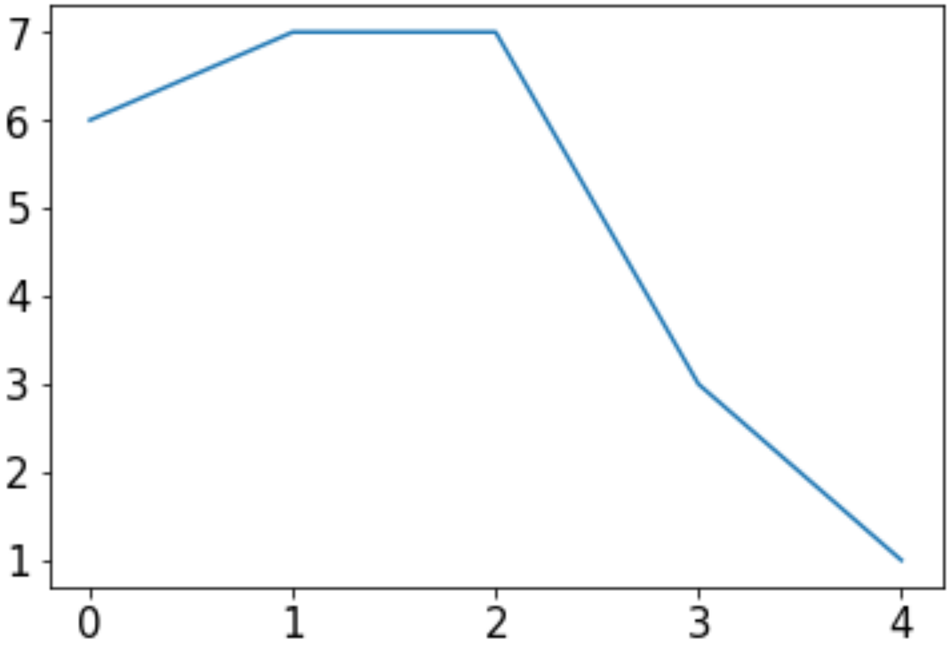

```
percents = []
for i in range(1, len(s)):
    diff = 100 * (s[i] / s[i-1] - 1)
    percents.append(diff)
Series(percents).plot.line()
```

what are we computing for diff?

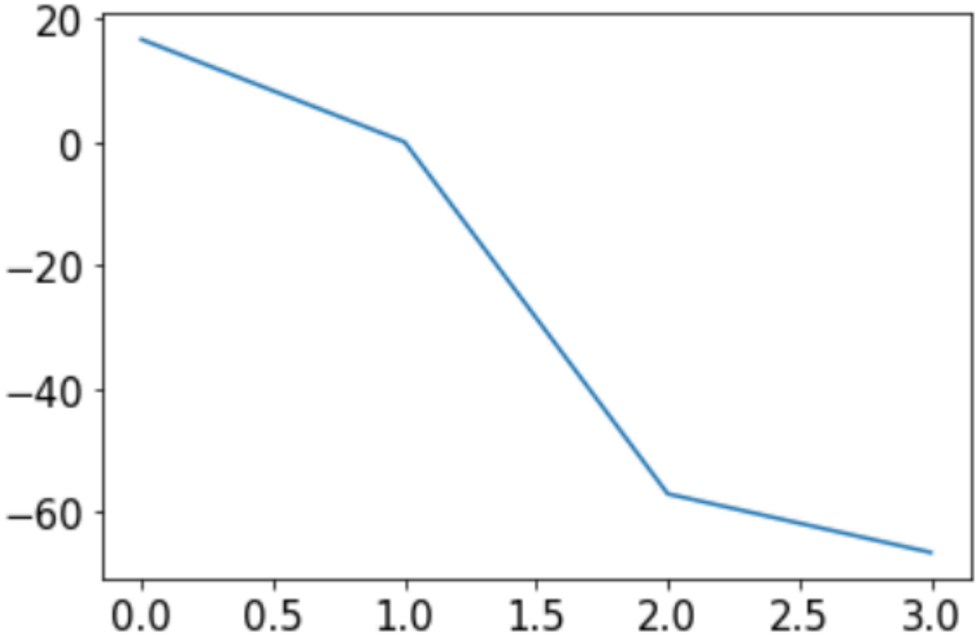

Example: change over time

```
s = Series(choice(10, size=5))
```

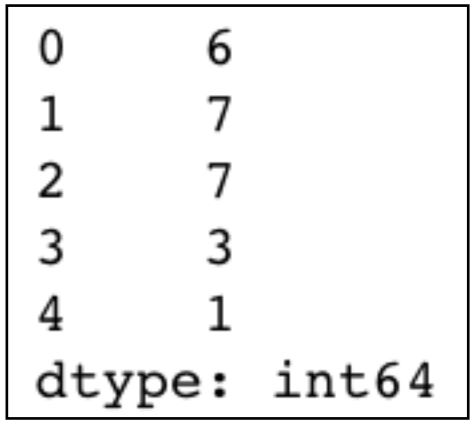

```
0      6
1      7
2      7
3      3
4      1
dtype: int64
```

```
s.plot.line()
```

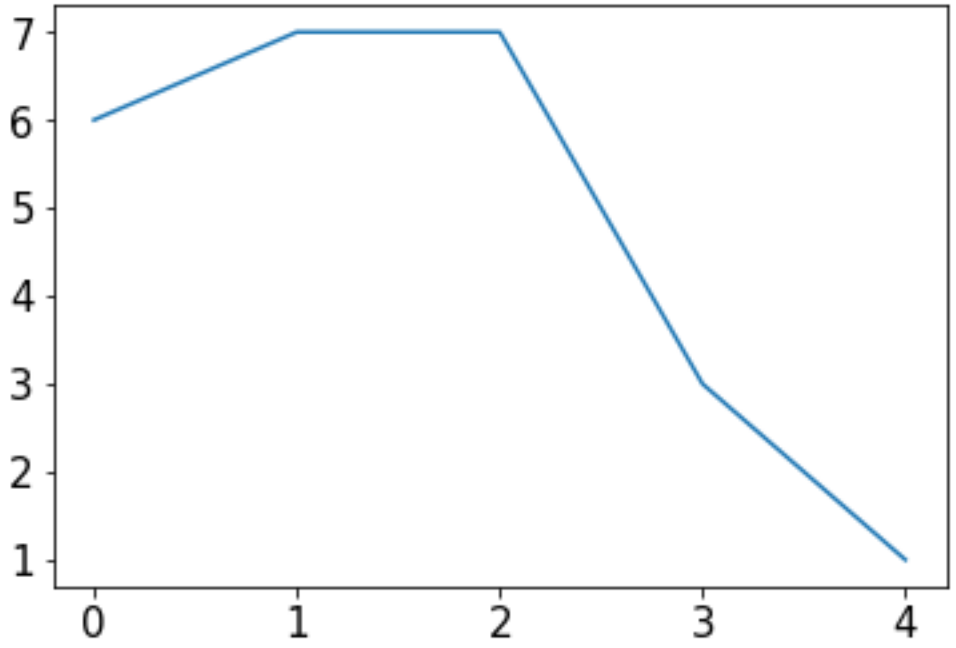

```
percents = []
for i in range(1, len(s)):
    diff = 100 * (s[i] / s[i-1] - 1)
    percents.append(diff)
Series(percents).plot.line()
```

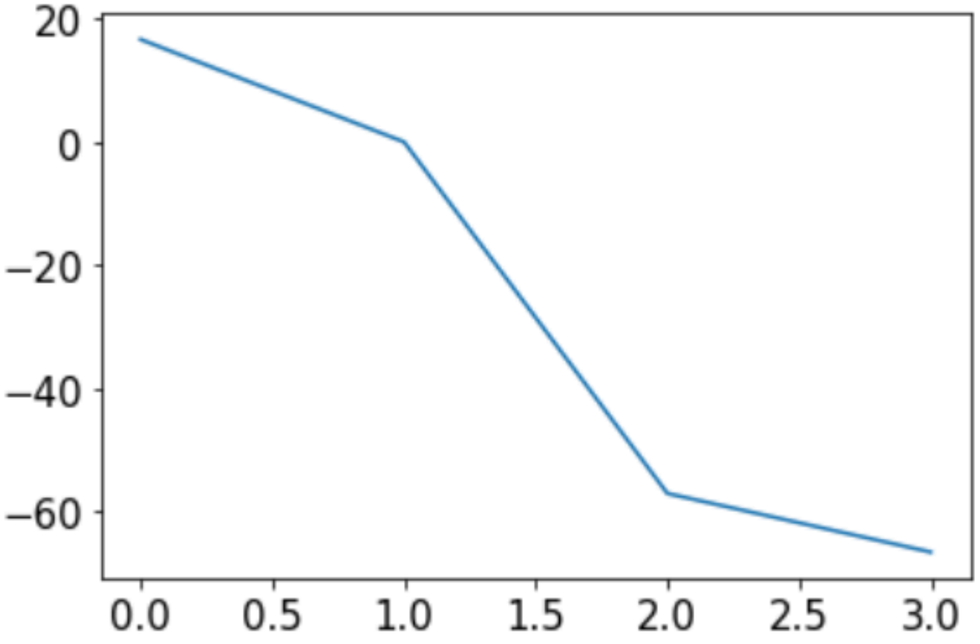

can you identify the bug in the code?

Example: change over time

```
s = Series(choice(10, size=5))
```

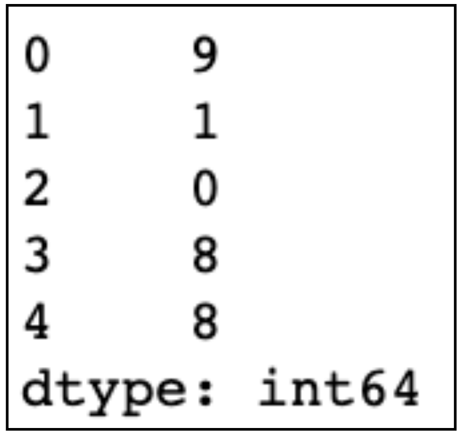

```
0    9
1    1
2    0
3    8
4    8
dtype: int64
```

```
s.plot.line()
```

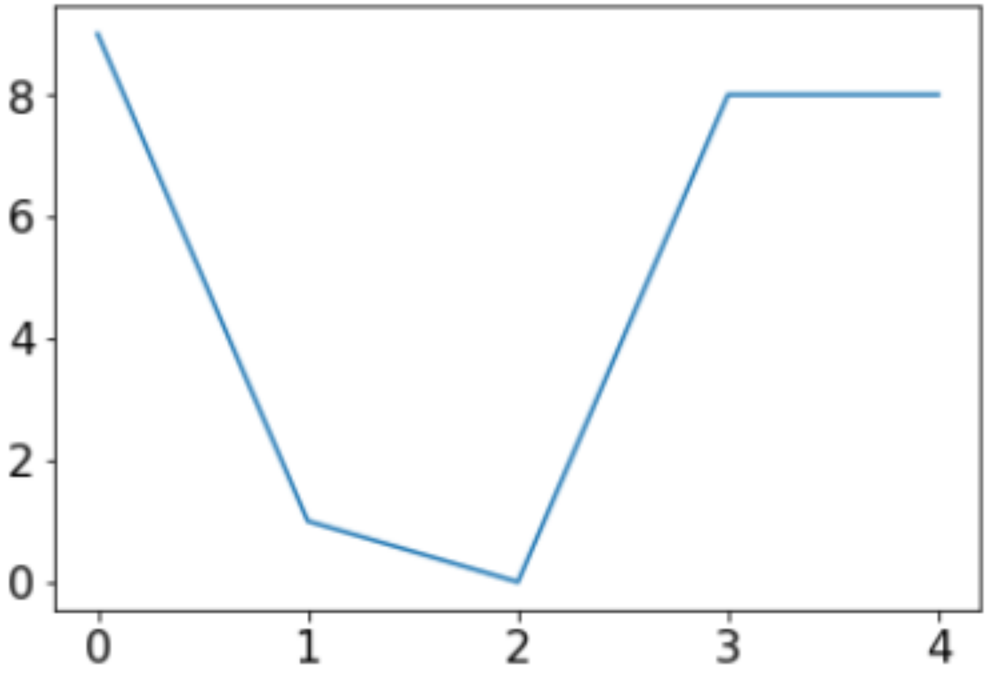

```
percents = []
for i in range(1, len(s)):
    diff = 100 * (s[i] / s[i-1] - 1)
    percents.append(diff)
Series(percents).plot.line()
```

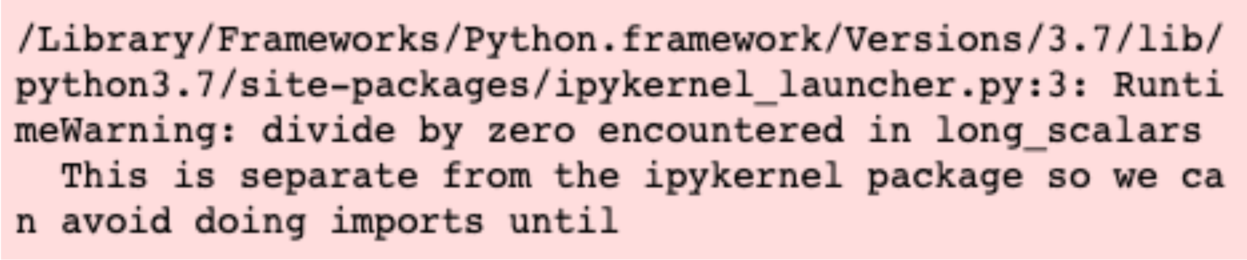

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/
python3.7/site-packages/ipykernel_launcher.py:3: Runti
meWarning: divide by zero encountered in long_scalars
  This is separate from the ipykernel package so we ca
n avoid doing imports until
```

can you identify the bug in the code?

Not all bugs are equal!

scary bugs

non-deterministic

Igor Siwanowicz

https://owlcation.com/stem/5-Badass-Bugs-That-You-Should-Have-Nightmares-About

"nice" bugs

deterministic (reproducible)

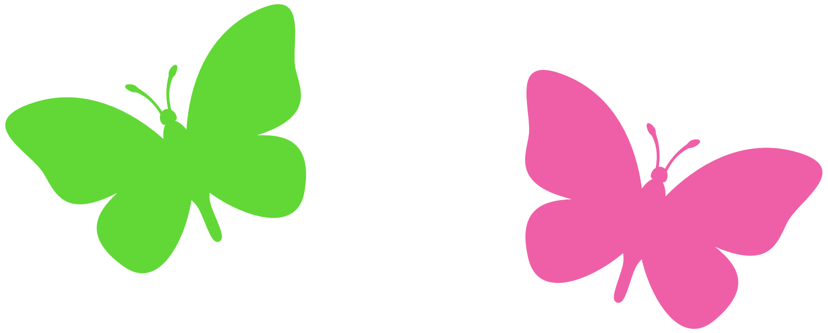

Not all bugs are equal!

scary bugs

non-deterministic
  system related
  randomness

Igor Siwanowicz

https://owlcation.com/stem/5-Badass-Bugs-That-You-Should-Have-Nightmares-About

"nice" bugs

deterministic (reproducible)

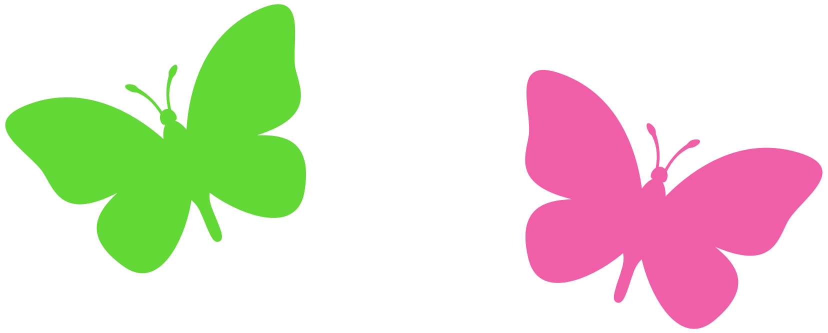

Not all bugs are equal!

scary bugs	"nice" bugs
non-deterministic system related randomness	**deterministic (reproducible)**
small data	**large data**
semantic	**syntax**

runtime

Igor Siwanowicz

https://owlcation.com/stem/5-Badass-Bugs-That-You-Should-Have-Nightmares-About

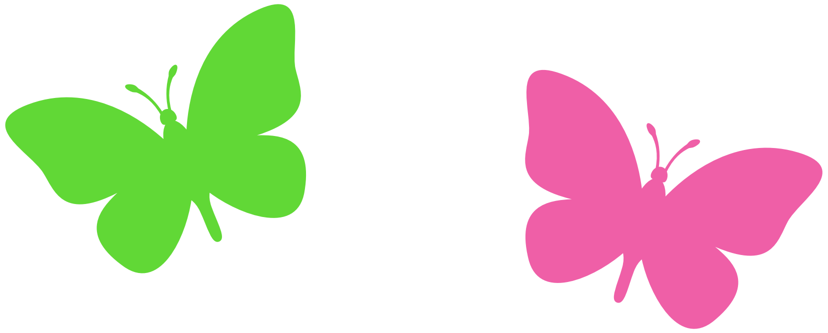

Not all bugs are equal!

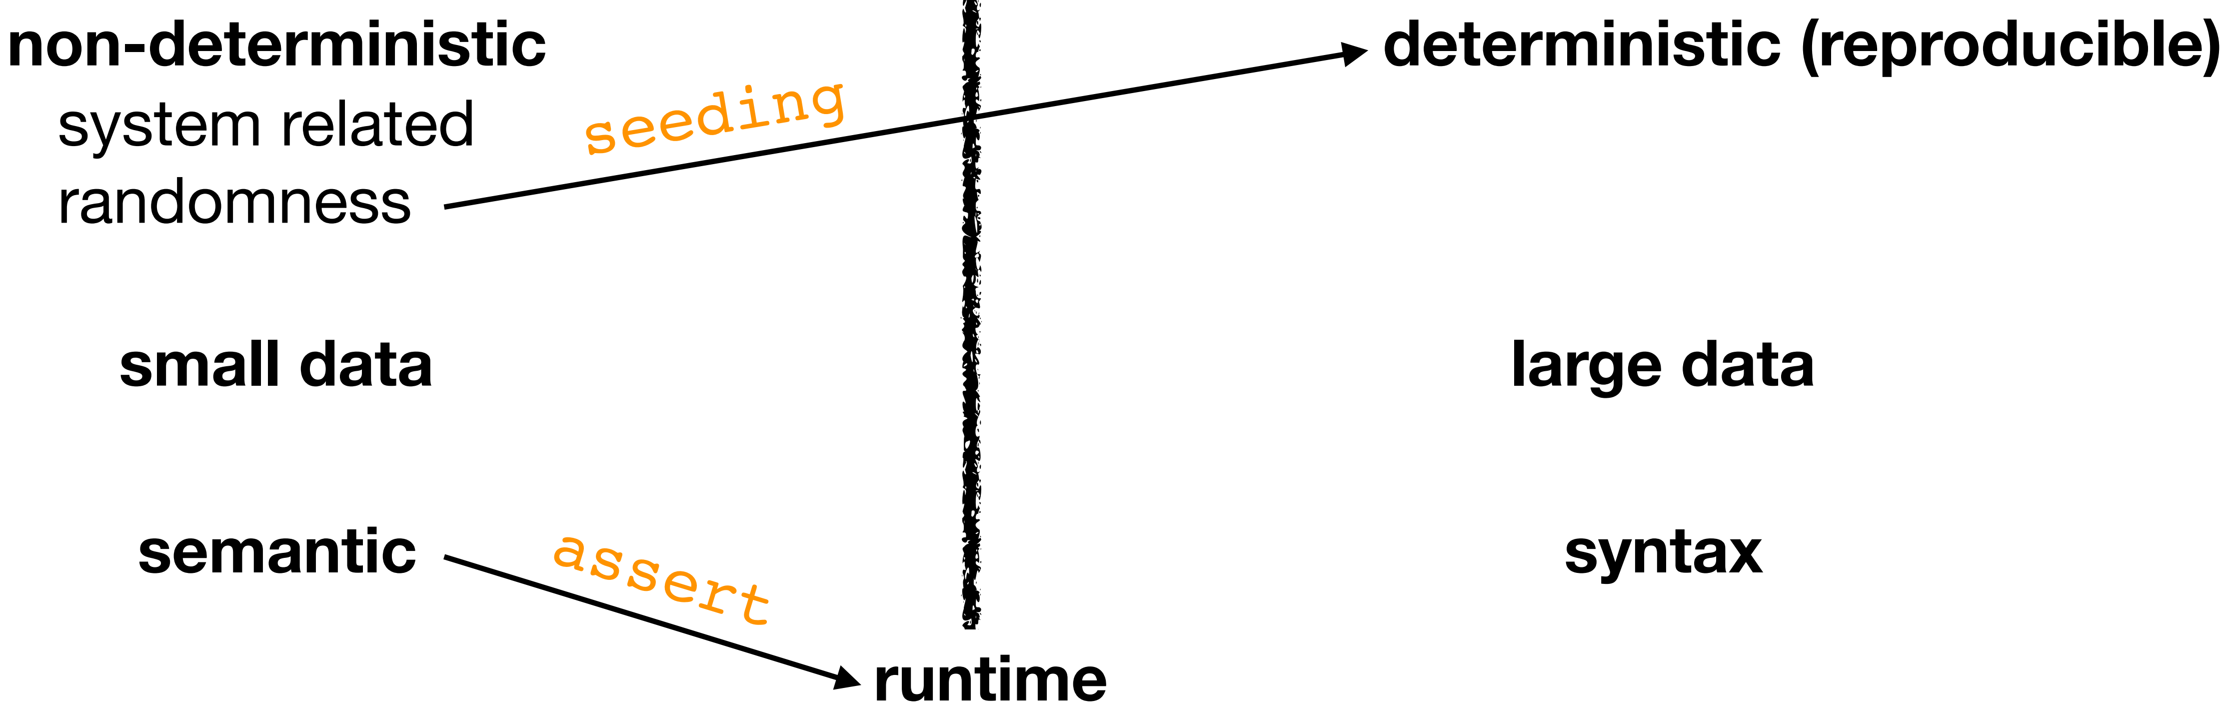

scary bugs

non-deterministic

system related

randomness

small data

semantic

runtime

"nice" bugs

deterministic (reproducible)

large data

syntax

Igor Siwanowicz

https://owlcation.com/stem/5-Badass-Bugs-That-You-Should-Have-Nightmares-About

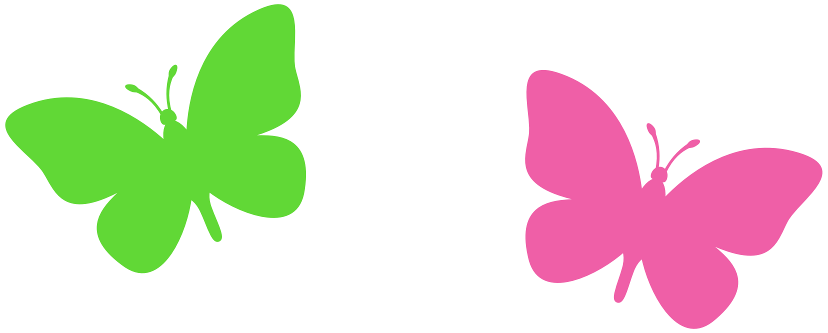

Not all bugs are equal!

scary bugs	"nice" bugs
non-deterministic system related randomness	**deterministic (reproducible)**
small data	**large data**
semantic	**syntax**

????

runtime

Igor Siwanowicz

https://owlcation.com/stem/5-Badass-Bugs-That-You-Should-Have-Nightmares-About

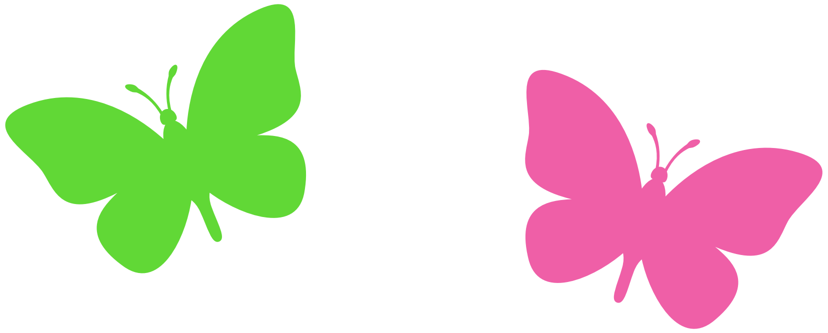

Pseudorandom Generators

"Random" generators are really just *pseudorandom*

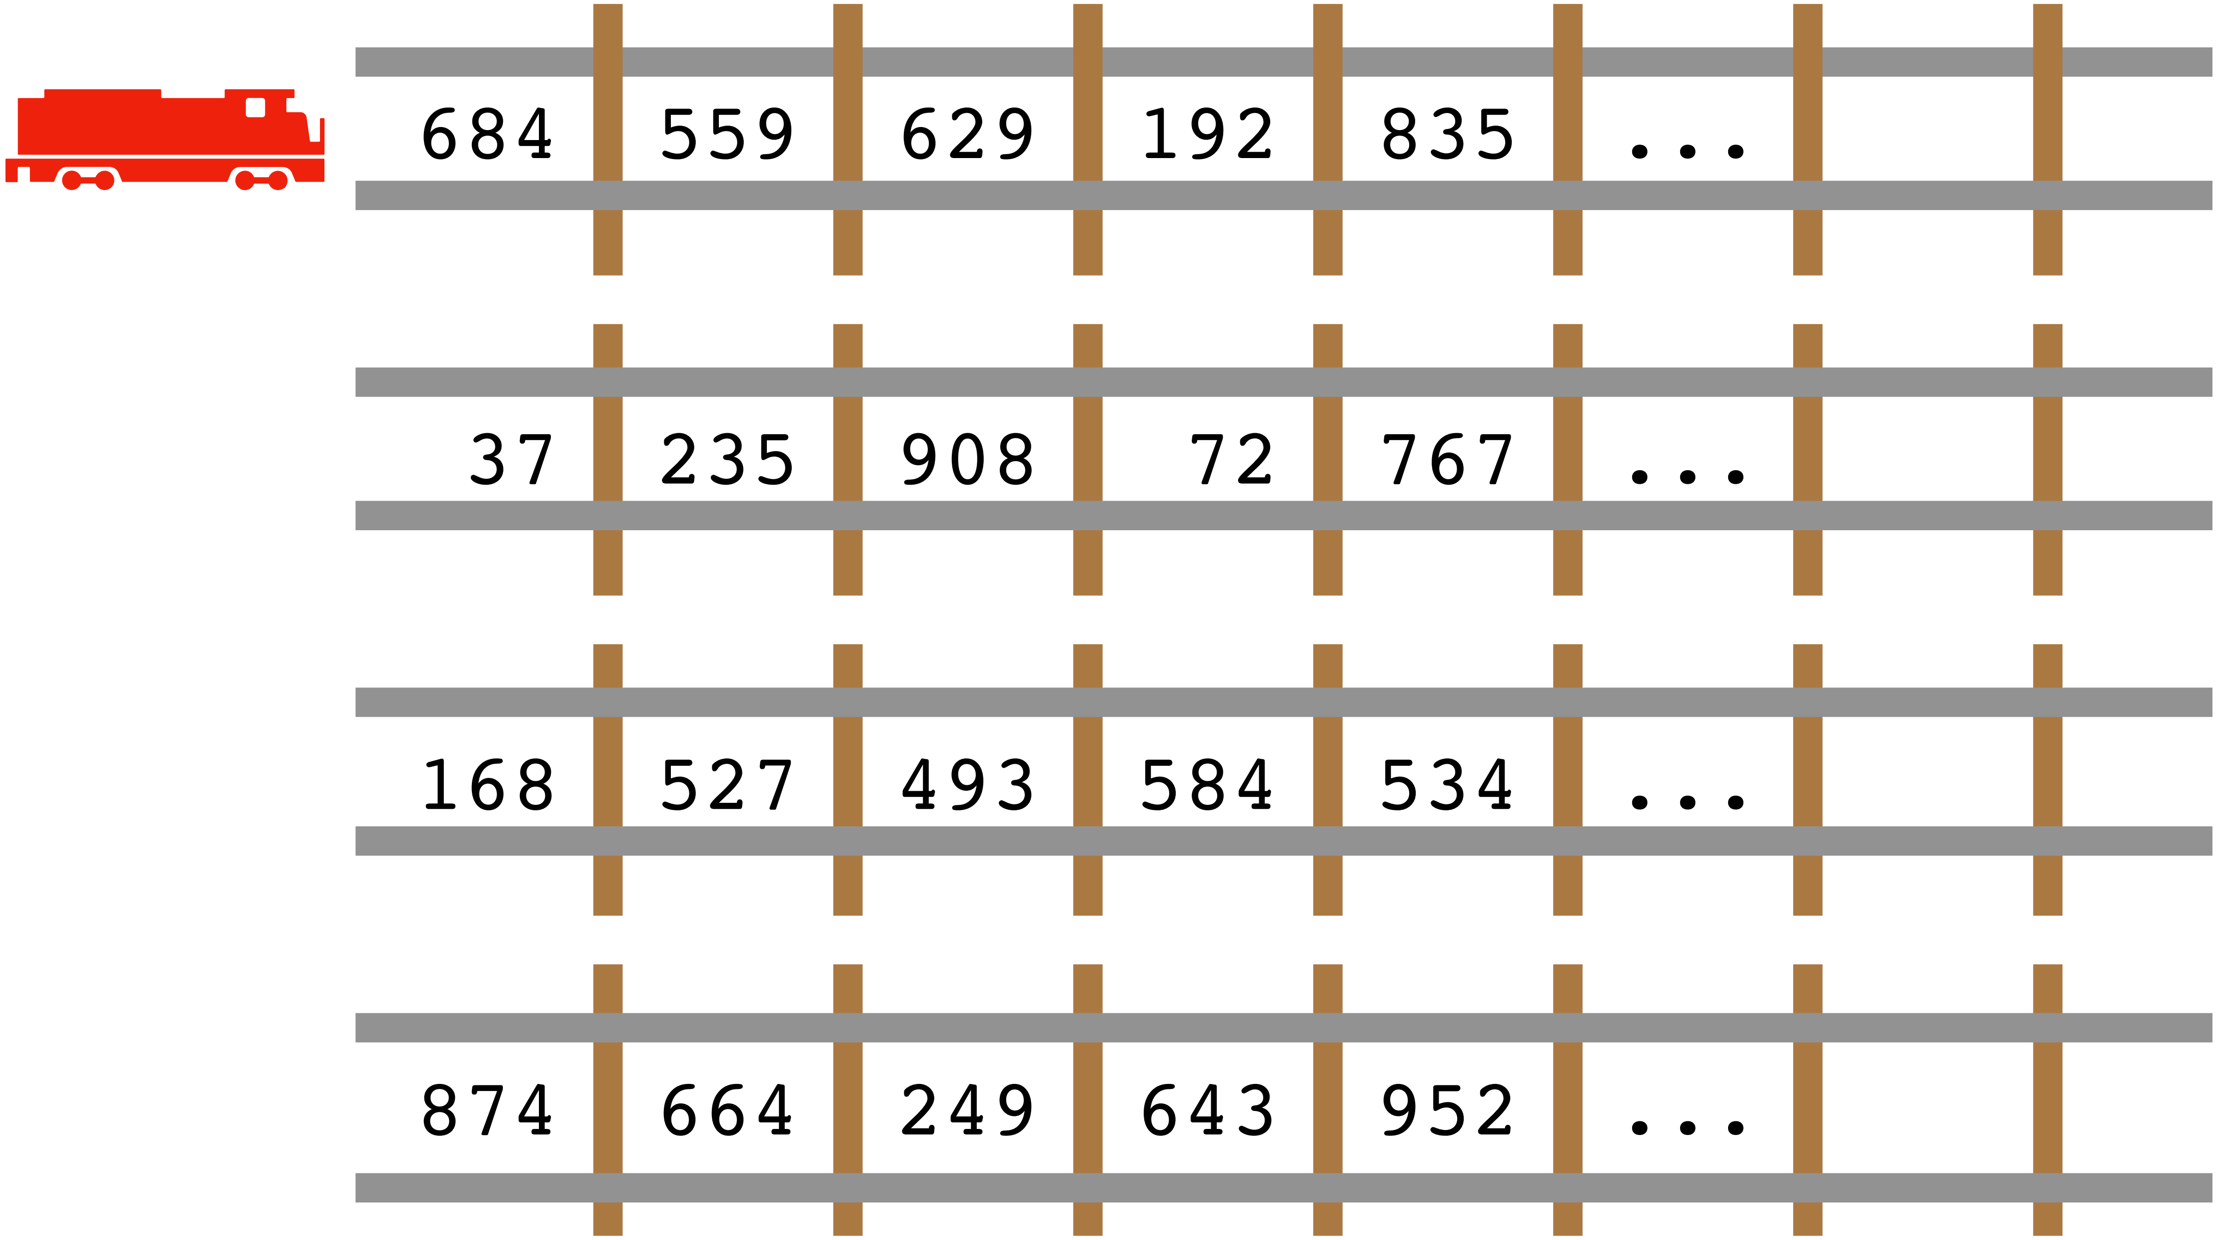

Pseudorandom Generators

Producing random numbers is like cruising down the tracks...

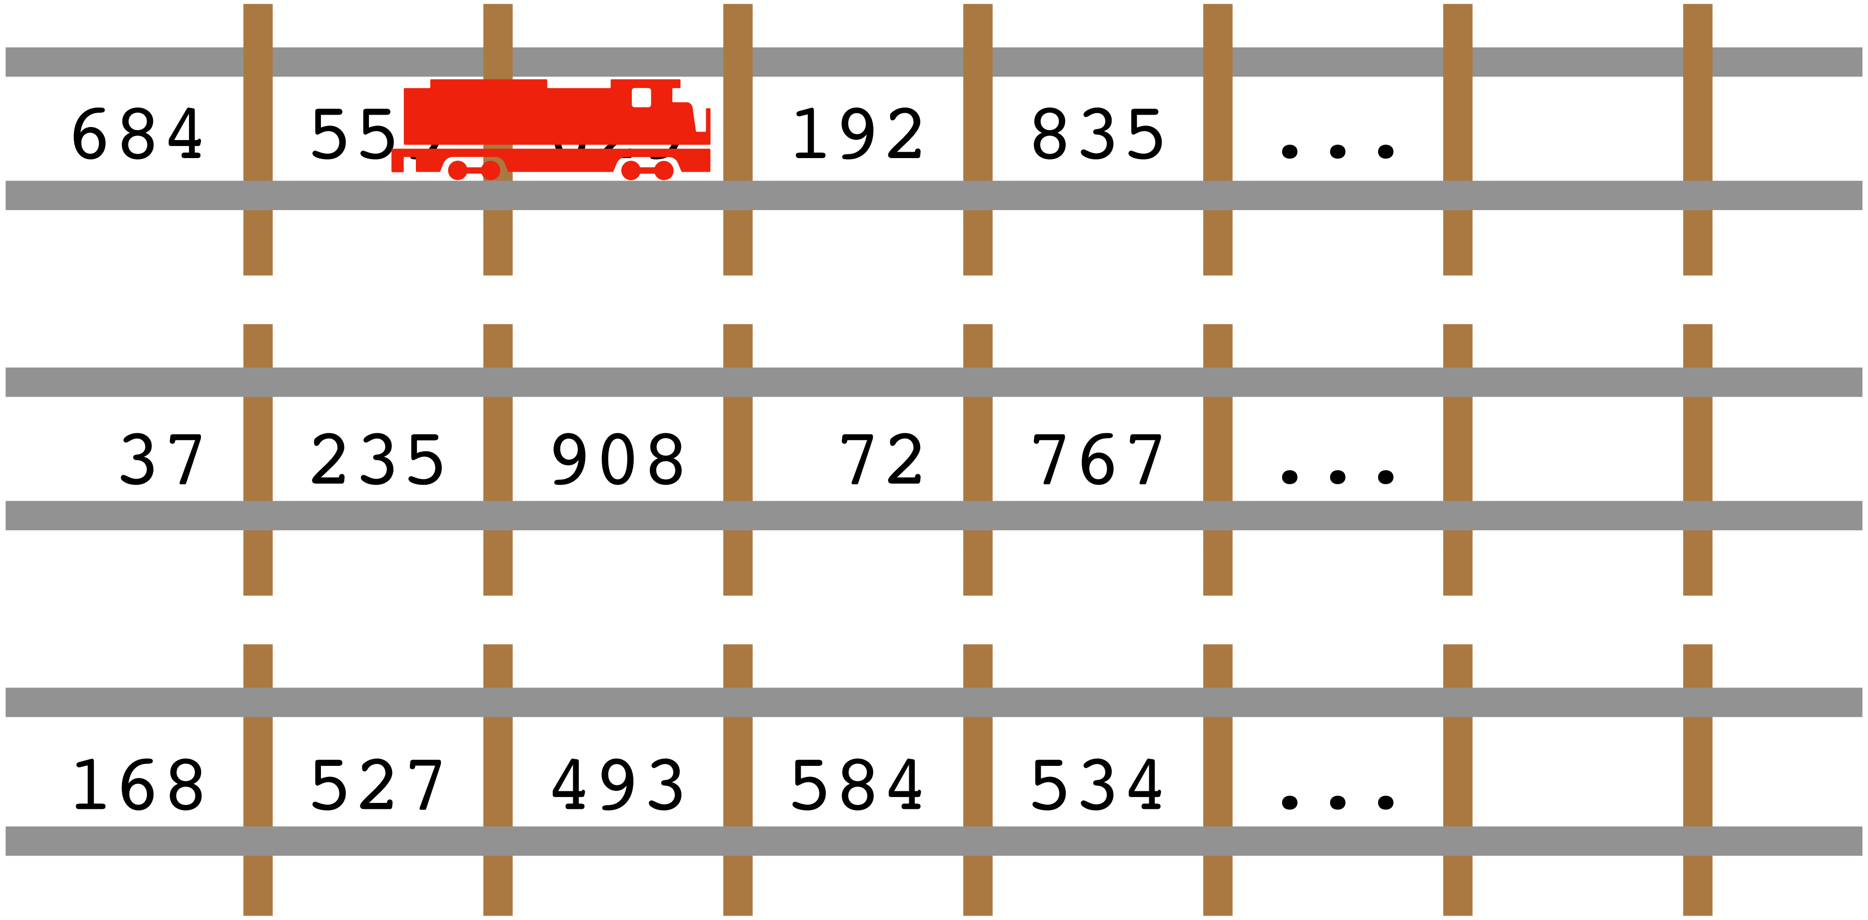

684	55		192	835	...
37	235	908	72	767	...
168	527	493	584	534	...

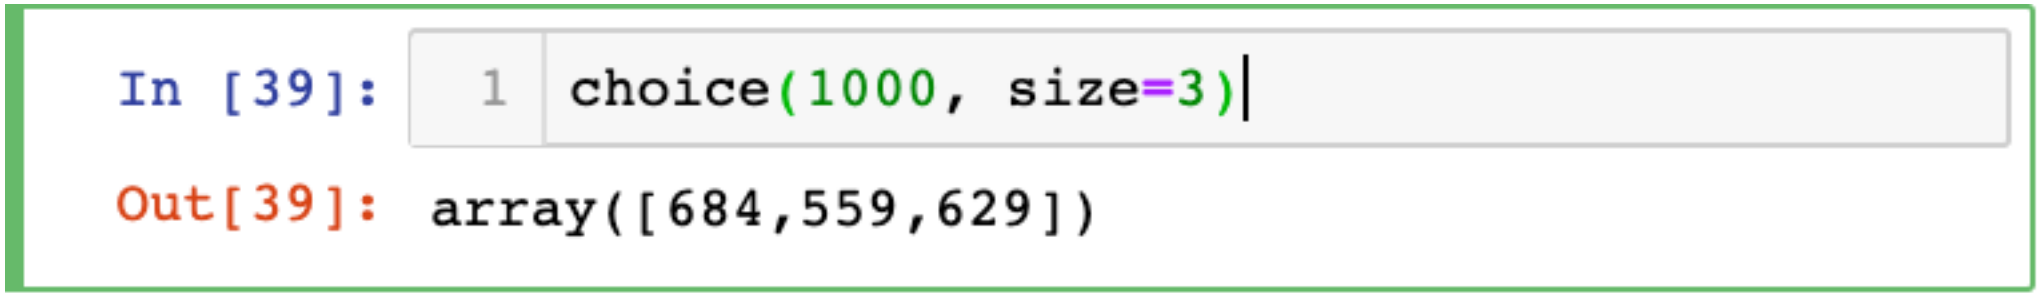

```
In [39]: choice(1000, size=3)
Out[39]: array([684,559,629])
```

Pseudorandom Generators

Every run, you get on another tracks, so it **feels** random

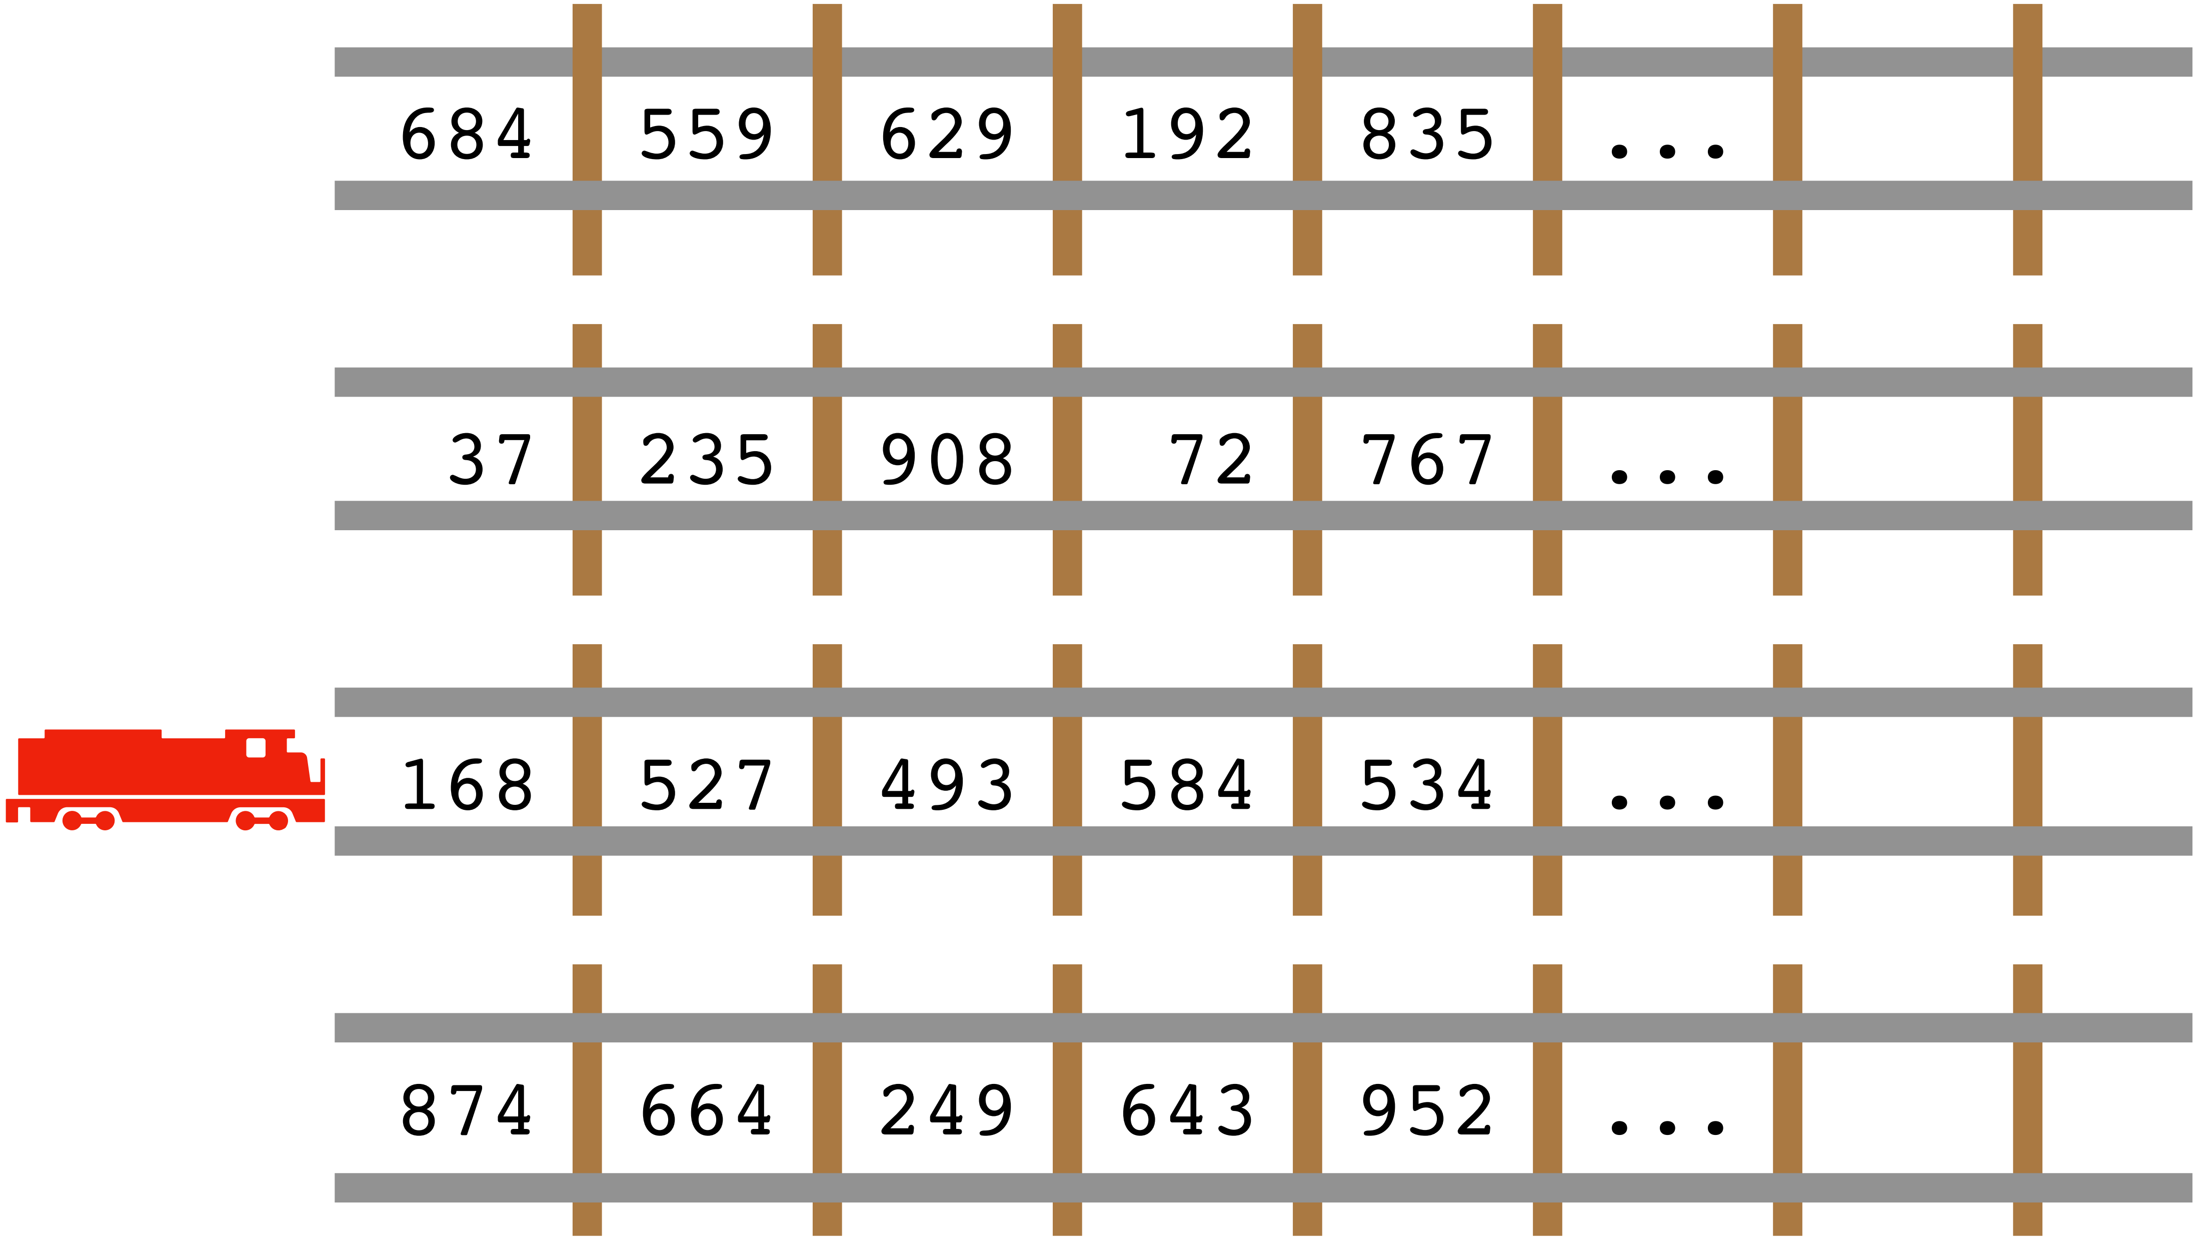

684	559	629	192	835	...
37	235	908	72	767	...
168	527	493	584	534	...
874	664	249	643	952	...

Seeding

What if I told you that you can **choose** your track?

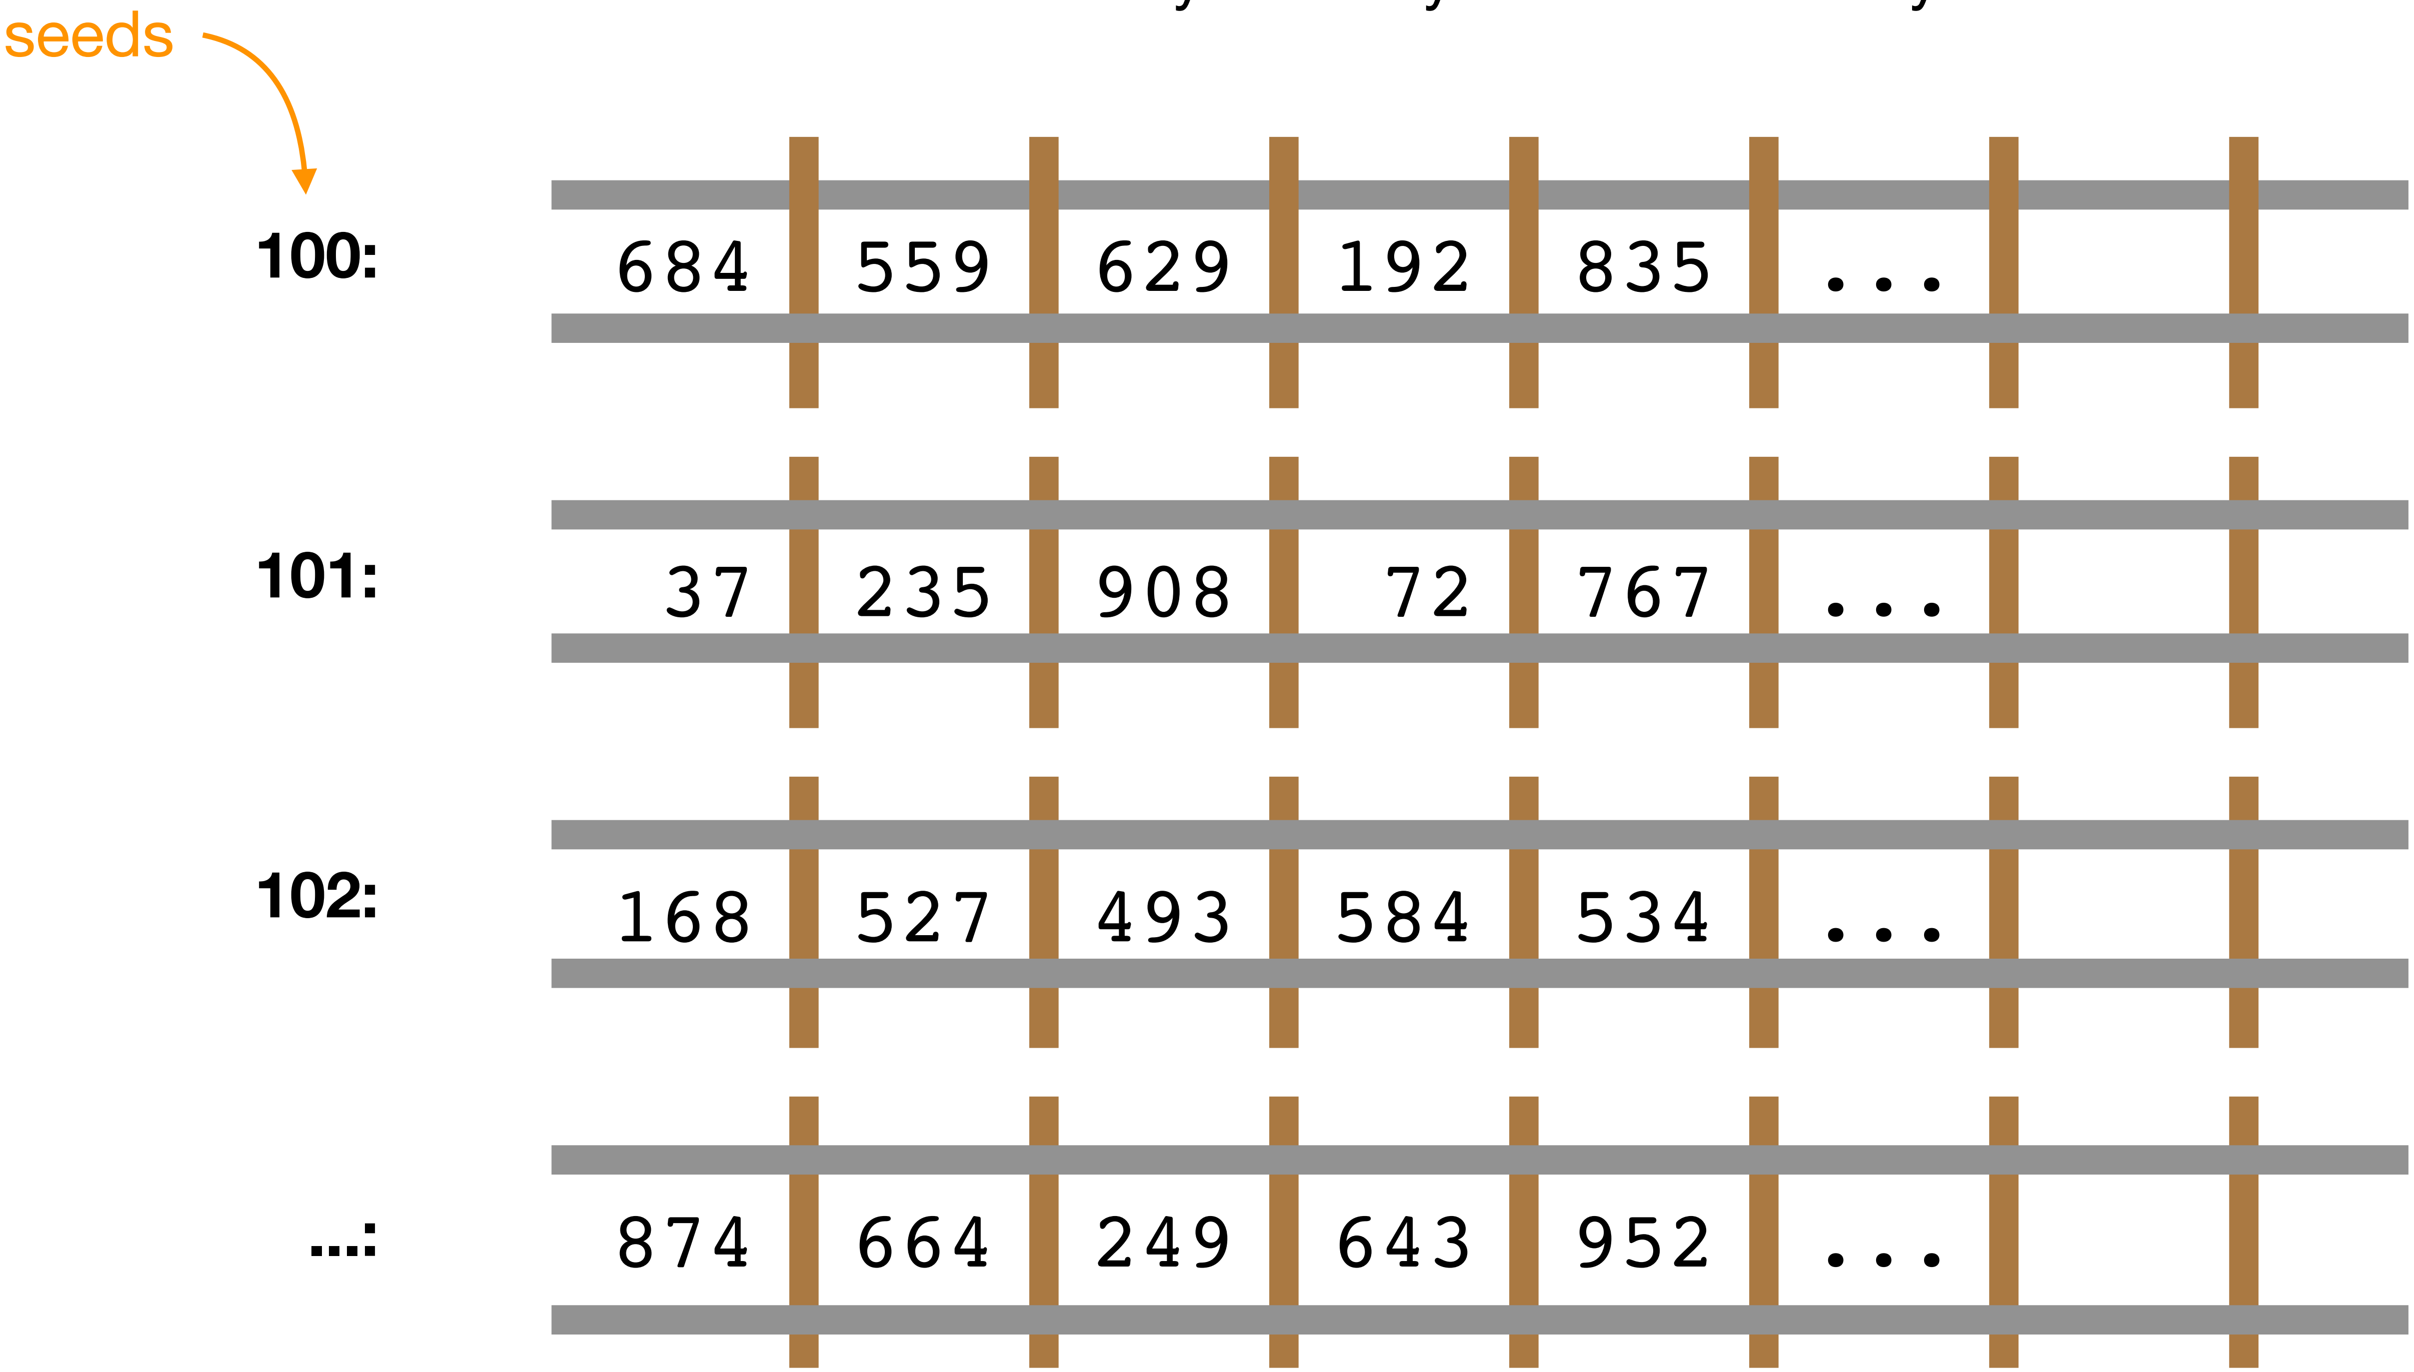

100:	684	559	629	192	835	...
101:	37	235	908	72	767	...
102:	168	527	493	584	534	...
...:	874	664	249	643	952	...

Seeding

What if I told you that you can **choose** your track?

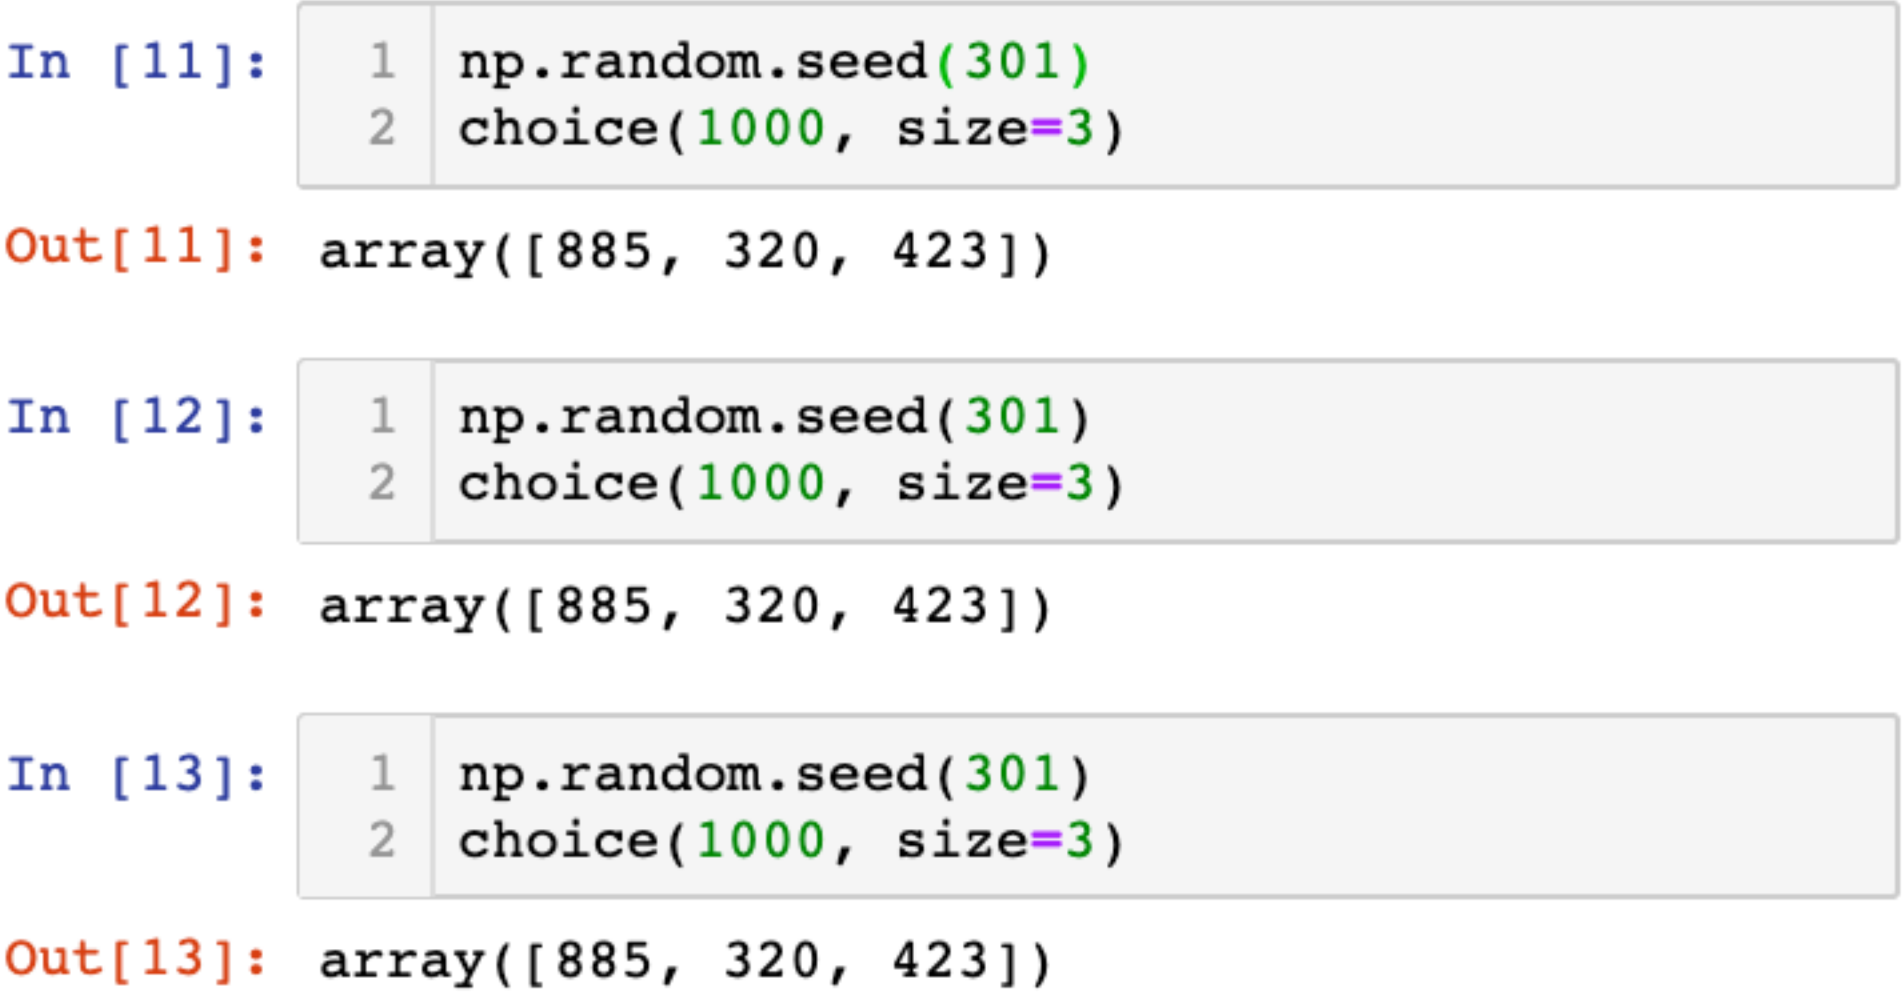

```
In [11]:  np.random.seed(301)
          choice(1000, size=3)

Out[11]: array([885, 320, 423])

In [12]:  np.random.seed(301)
          choice(1000, size=3)

Out[12]: array([885, 320, 423])

In [13]:  np.random.seed(301)
          choice(1000, size=3)

Out[13]: array([885, 320, 423])
```

Seeding

Common approach for simulations:

1. seed using current time
2. print seed
3. use the seed for reproducing bugs, as necessary

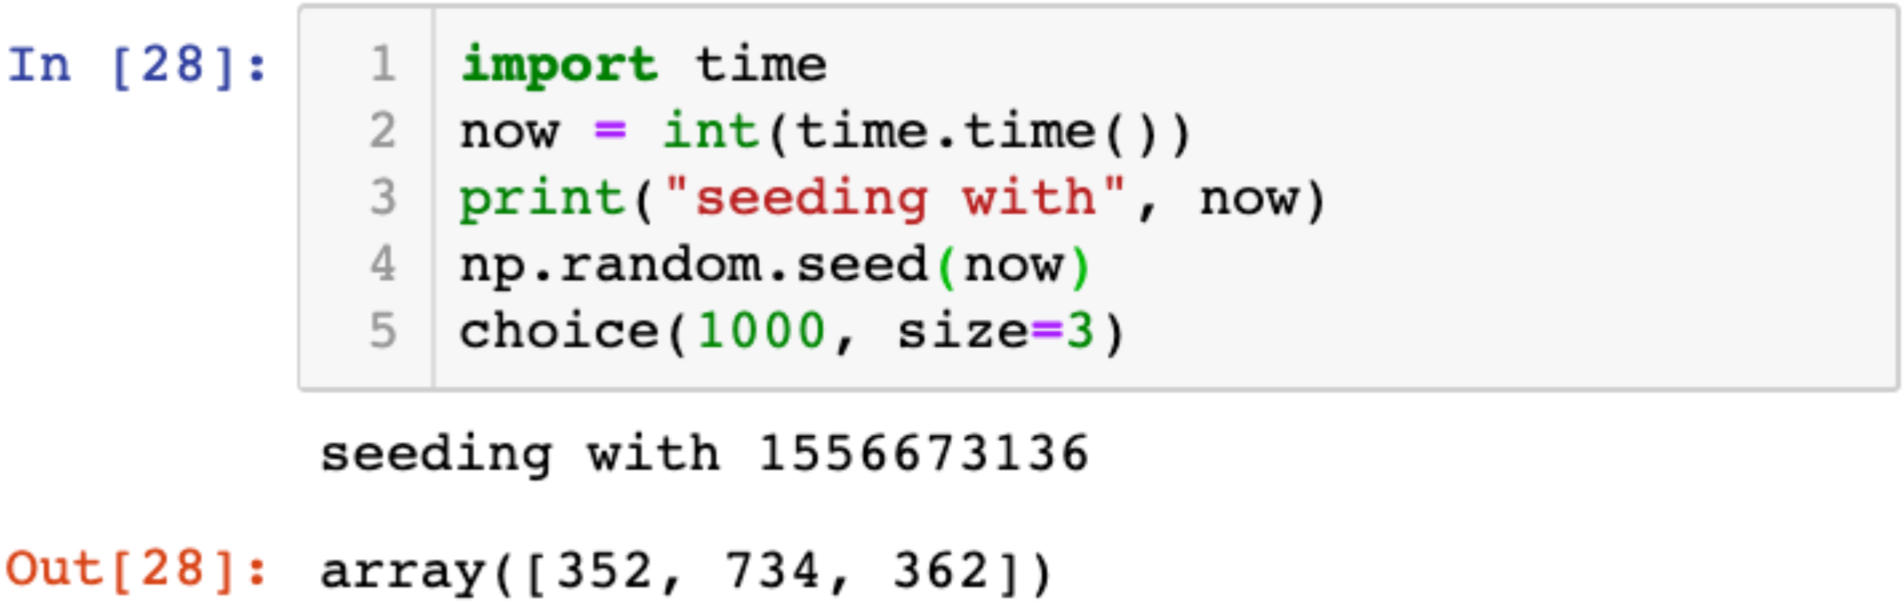

```
In [28]: import time
         now = int(time.time())
         print("seeding with", now)
         np.random.seed(now)
         choice(1000, size=3)

         seeding with 1556673136

Out[28]: array([352, 734, 362])
```

Outline

choice()

bugs and seeding

significance

histograms

normal()

In a noisy world, what is noteworthy?

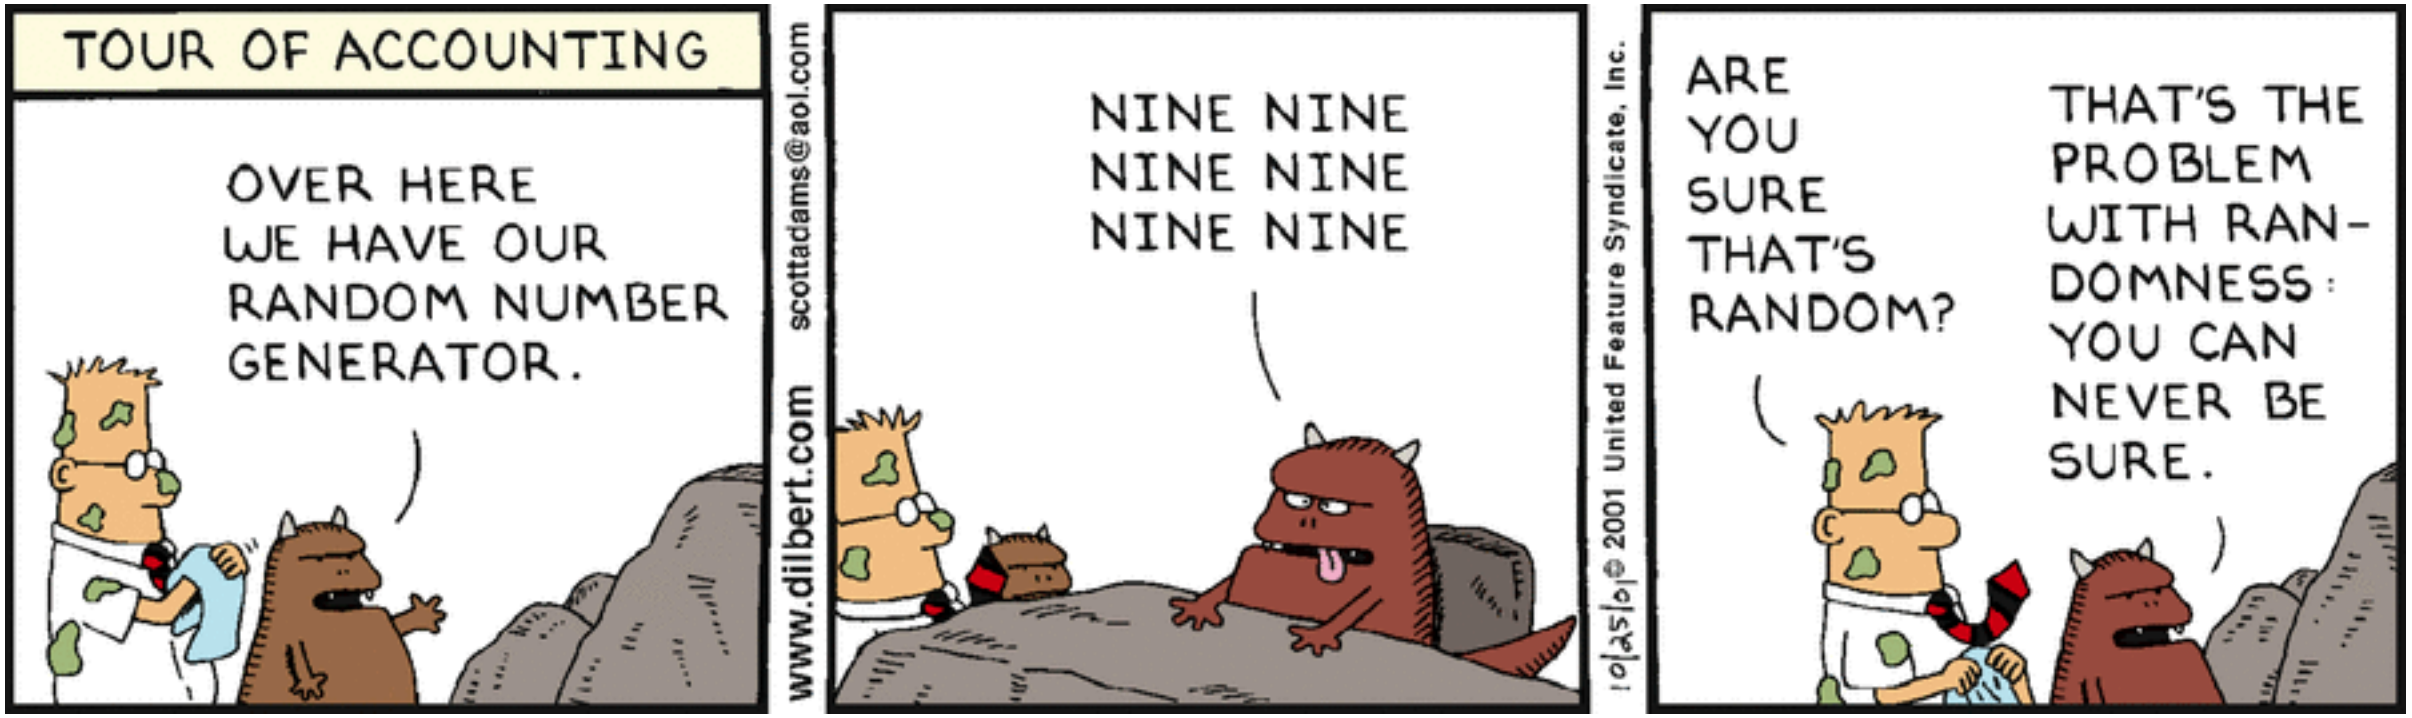

https://dilbert.com/strip/2001-10-25

Is this coin biased?

51

49

Call shenanigans?

whoever has the coin cheated
(it's not 50/50 heads/tails)

Is this coin biased?

51 **49**

Call shenanigans?

whoever has the coin cheated
(it's not 50/50 heads/tails)

a statistician might say we're trying to decide if the evidence that the coin isn't fair is **statistically significant**

Is this coin biased?

Call shenanigans? No.

Is this coin biased?

51

49

Call shenanigans? No.

5

95

Call shenanigans?

Is this coin biased?

51 49

Call shenanigans? No.

5 95

Call shenanigans? Yes.

Note: there is a non-zero probability that a fair coin will do this, but the odds are slim

Is this coin biased?

51 **49**

Call shenanigans? No.

5 **95**

Call shenanigans? Yes.

55 **45**

Call shenanigans?

55 million **45 million**

Call shenanigans?

Is this coin biased?

Call shenanigans? No.

Call shenanigans? Yes.

Call shenanigans? No.

Call shenanigans? Yes.

Is this coin biased?

51 **49**

Call shenanigans? No.

5 **95**

large skew is good evidence of shenanigans

Call shenanigans? Yes.

55 **45**

Call shenanigans? No.

55 million **45 million**

small skew over **large samples** is good evidence

Call shenanigans? Yes.

Demo: CoinSim

60

40

Call shenanigans?

Strategy: simulate a fair coin

Demo: CoinSim

60

40

Call shenanigans?

Strategy: simulate a fair coin

1. "flip" it 100 times using numpy.random.choice
2. count heads
3. repeat above 10K times

```
[50, 61, 51, 44, 39, 43, 51, 49, 49, 38, ...]
```

Demo: CoinSim

60

40

Call shenanigans?

we got 10 more heads than we expect on average

Strategy: simulate a fair coin

1. "flip" it 100 times using numpy.random.choice
2. count heads
3. repeat above 10K times

```
[50, 61, 51, 44, 39, 43, 51, 49, 49, 38, ...]
```

Demo: CoinSim

60 **40**

Call shenanigans?

we got 10 more heads than we expect on average
how common is this?

Strategy: simulate a fair coin

1. "flip" it 100 times using numpy.random.choice
2. count heads
3. repeat above 10K times

```
[50, 61, 51, 44, 39, 43, 51, 49, 49, 38, ...]
```

Demo: CoinSim

60

40

Call shenanigans?

we got 10 more heads than we expect on average
how common is this?

Strategy: simulate a fair coin

1. "flip" it 100 times using numpy.random.choice
2. count heads
3. repeat above 10K times

```
[50, 61, 51, 44, 39, 43, 51, 49, 49, 38, ...]
```

11 more

12 less

Outline

choice()

bugs and seeding

significance

histograms

normal()

Frequencies across categories

bars are a good way to view frequencies across categories

```
s = Series(["rock", "rock", "paper",
            "scissors", "scissors", "scissors"])

s.value_counts().plot.bar(color="orange")
```

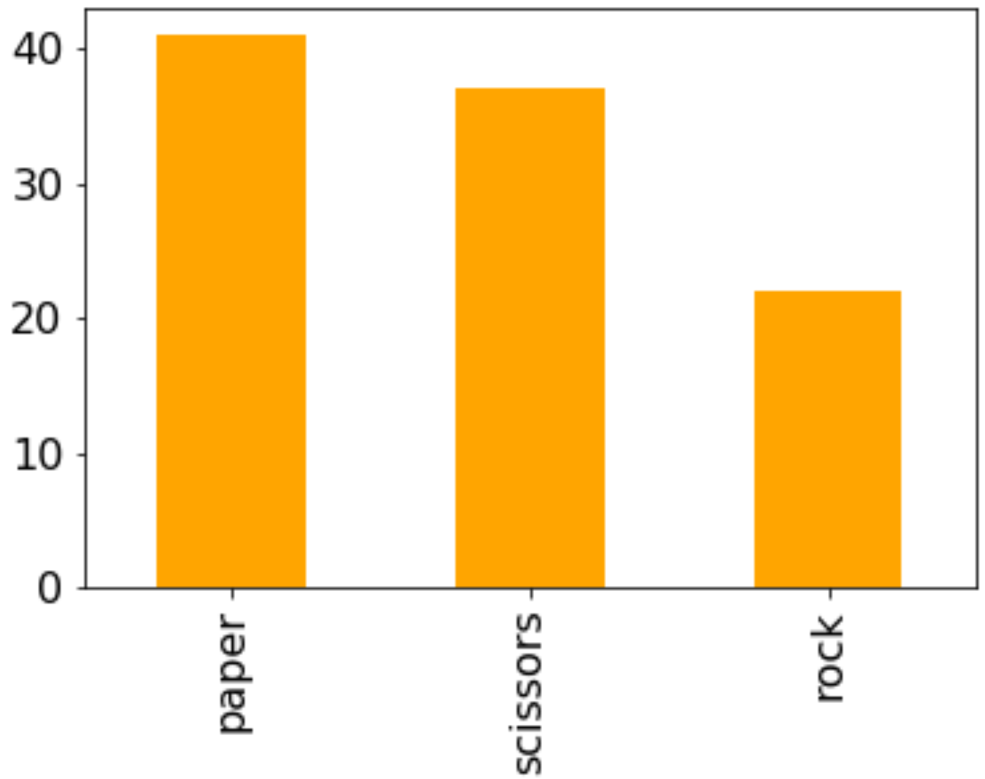

Frequencies across numbers

bars are a bad way to view frequencies across numbers

```
s = Series([0, 0, 1, 8, 9, 9])

s.value_counts().plot.bar(color="orange")
```

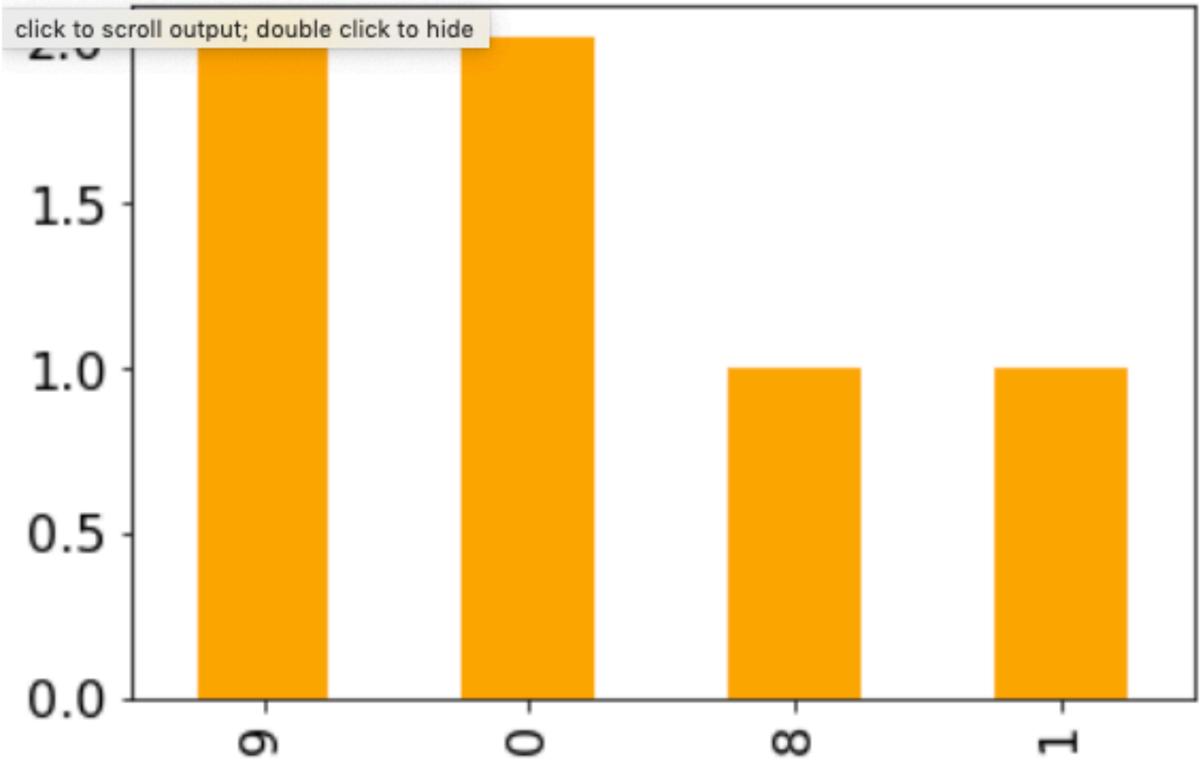

numbers not ordered

Frequencies across numbers

bars are a bad way to view frequencies across numbers

```
s = Series([0, 0, 1, 8, 9, 9])

s.value_counts().sort_index().plot.bar(color="orange")
```

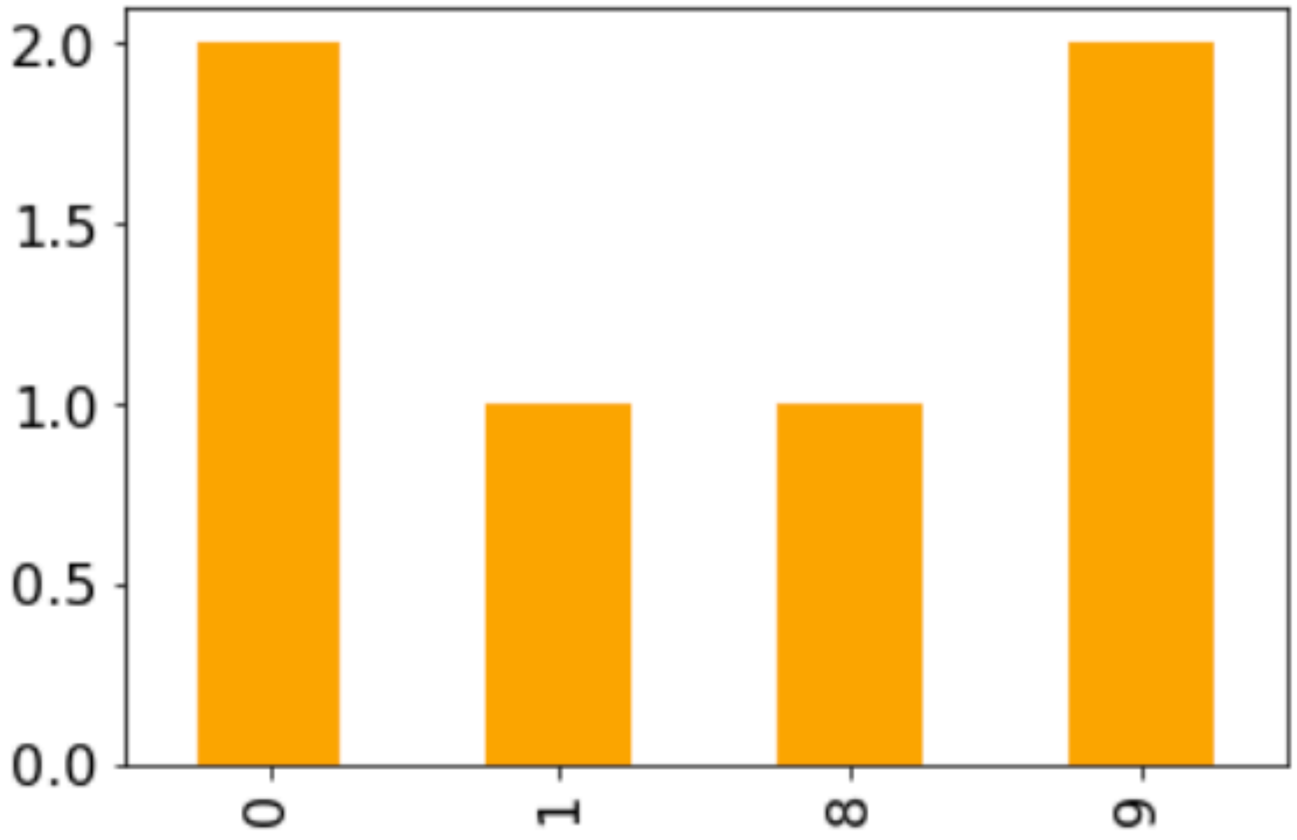

gap between 1 and 8 not obvious

Frequencies across numbers

bars are a bad way to view frequencies across numbers

```
s = Series([0, 0, 1, 8, 9, 9])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist()
```

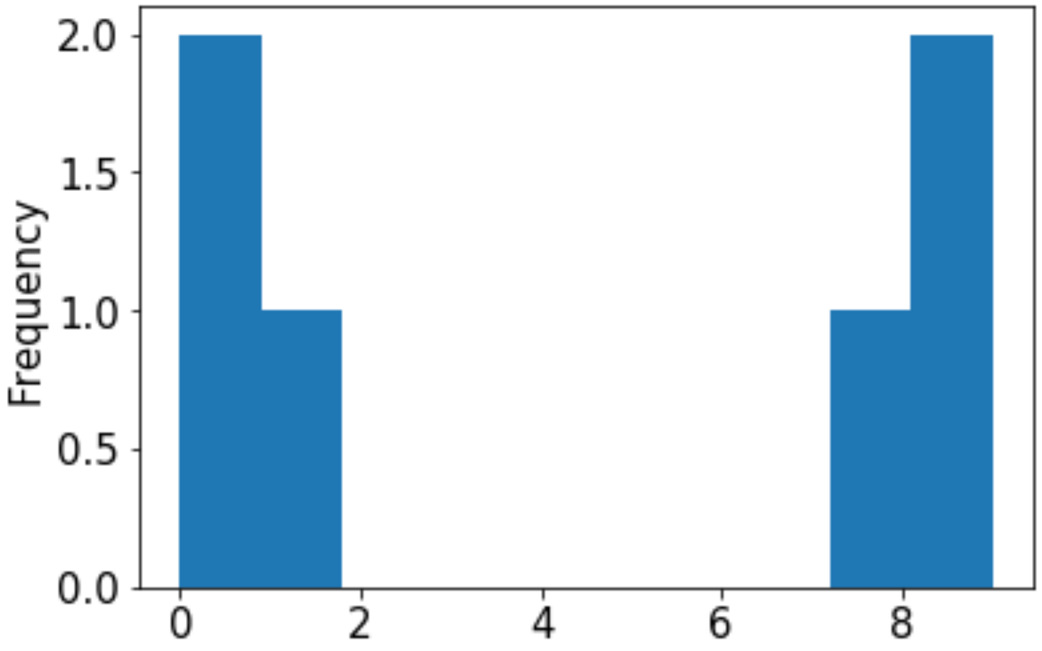

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0, 0, 1, 8, 9, 9])
```

```
~~s.value_counts().sort_index().plot.bar()~~
s.plot.hist()
```

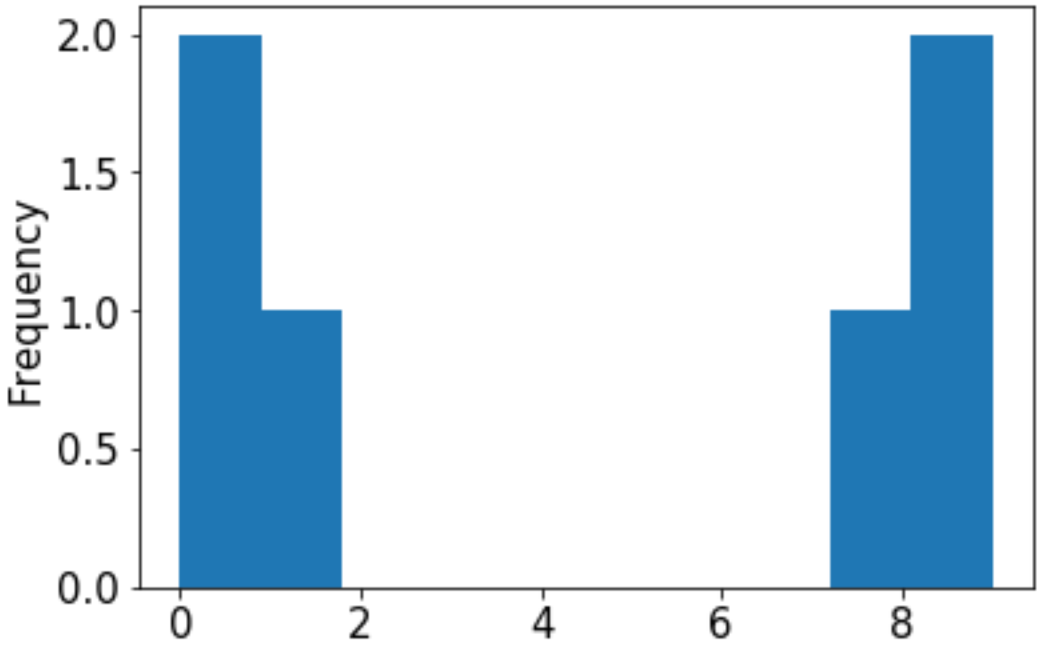

this kind of plot is called a histogram

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist()
```

both 0 and 0.1

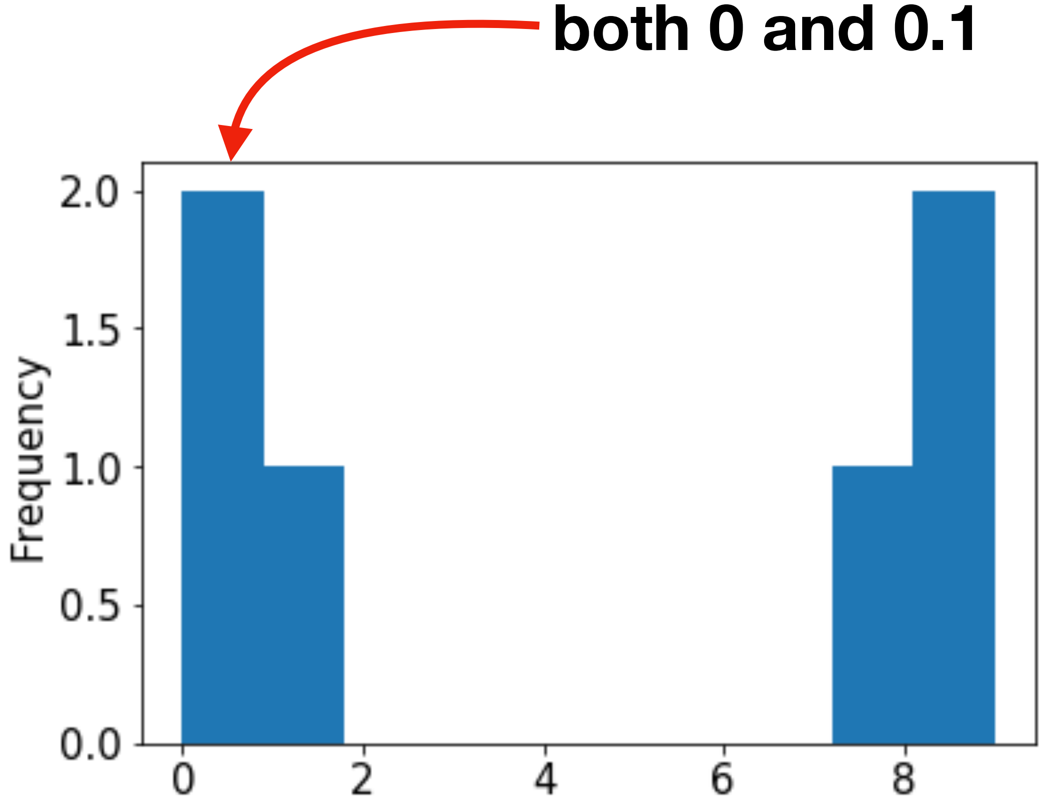

a histogram "bins" nearby numbers to create discrete bars

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist(bins=10)
```

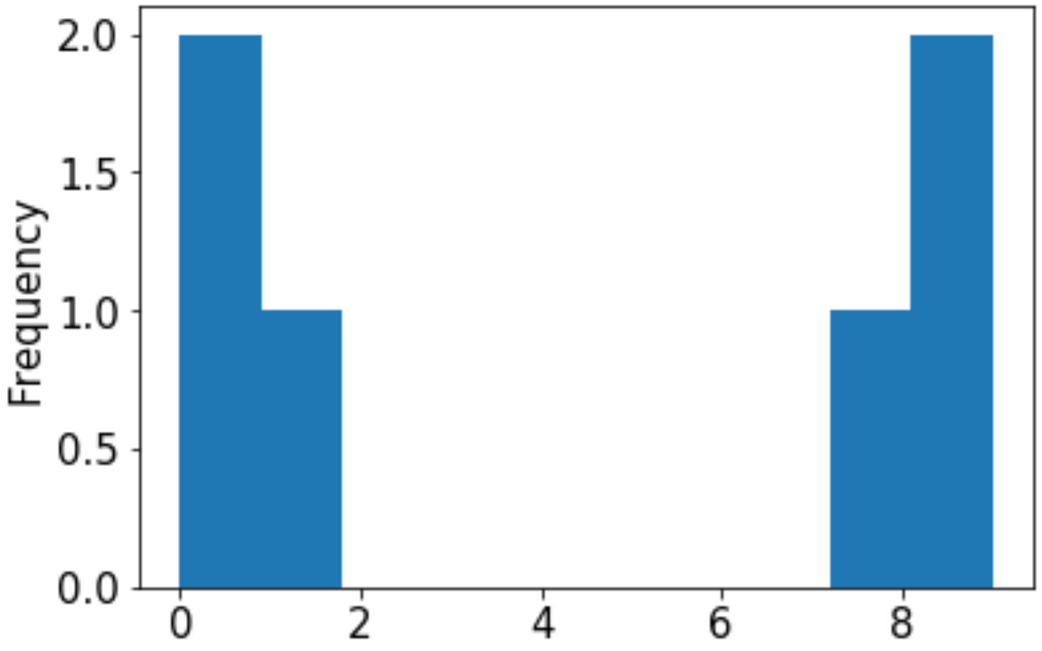

we can control the number of bins

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist(bins=3)
```

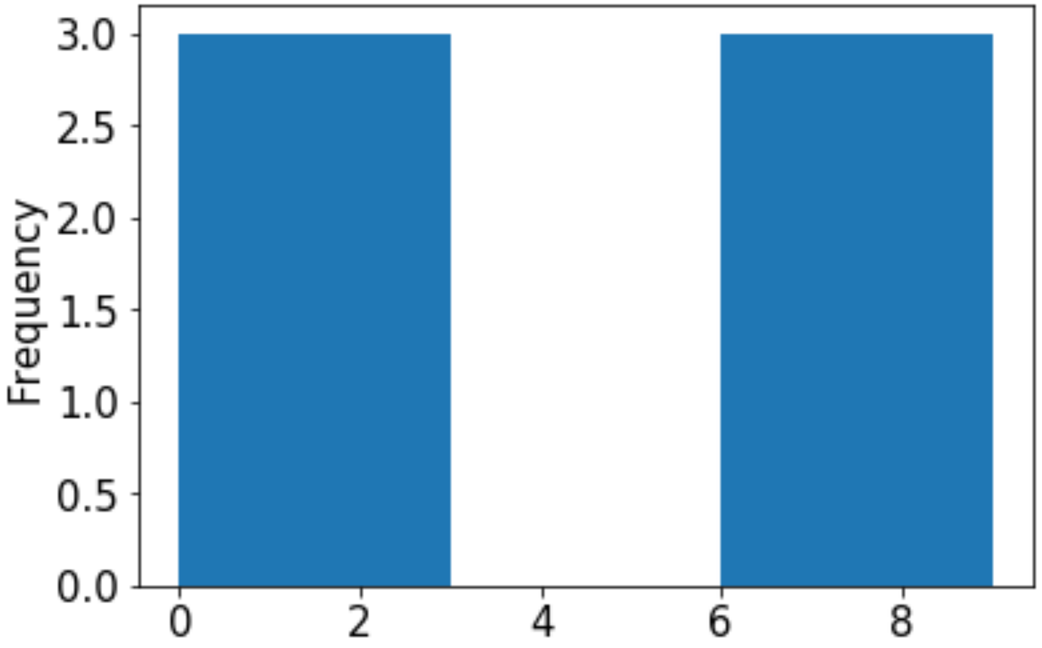

too few bins provides too little detail

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist(bins=100)
```

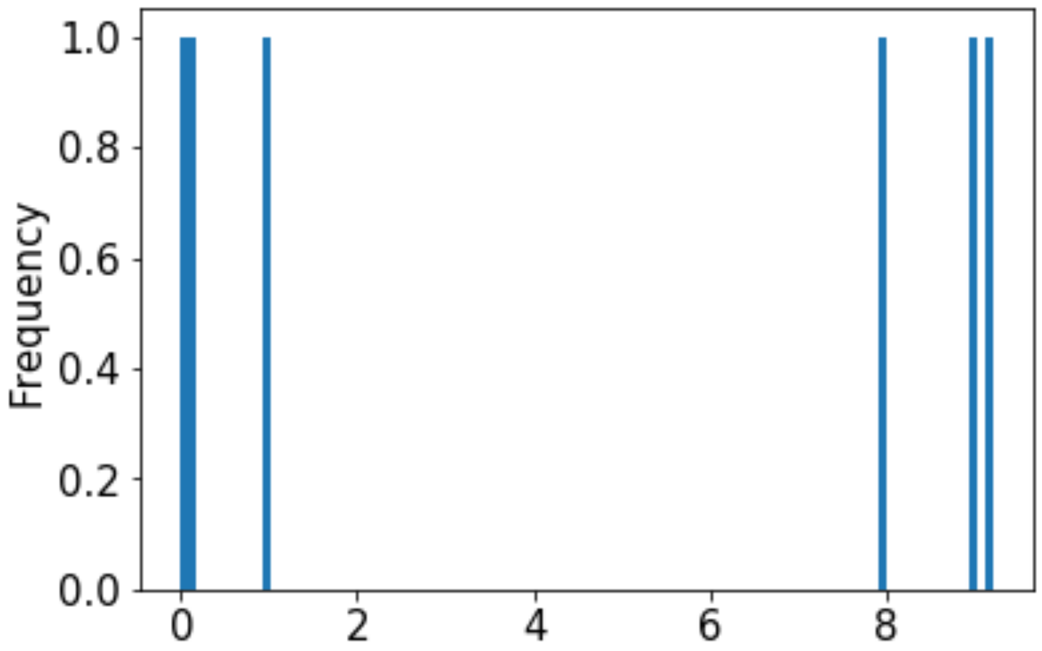

too many bins provides too much detail (equally bad)

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist(bins=10)
```

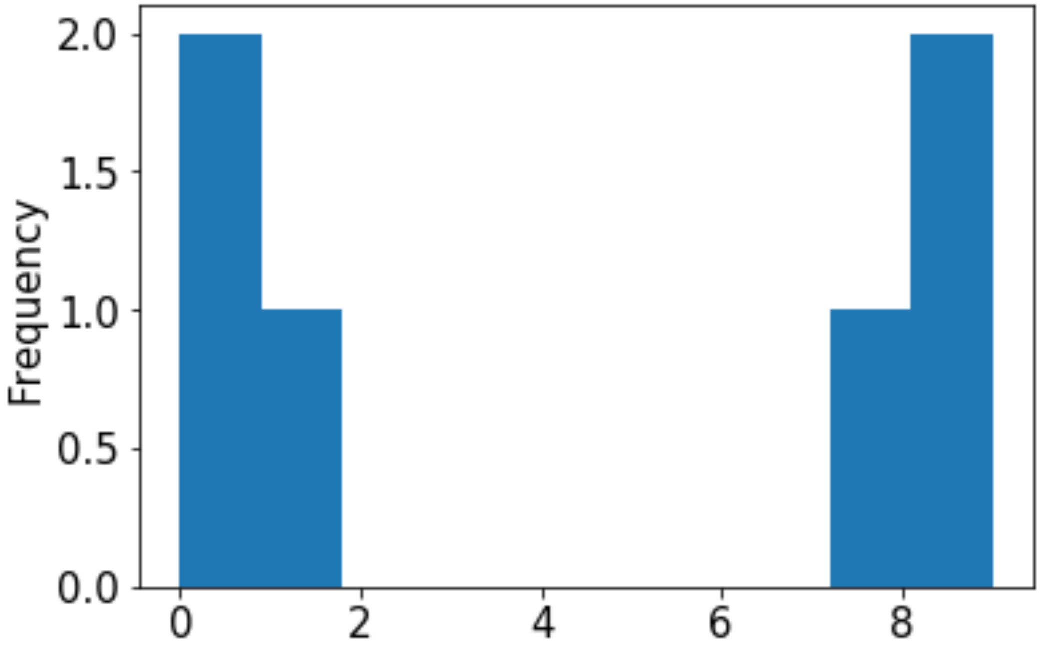

numpy chooses the default bin boundaries

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist(bins=[0,1,2,3,4,5,6,7,8,9,10])
```

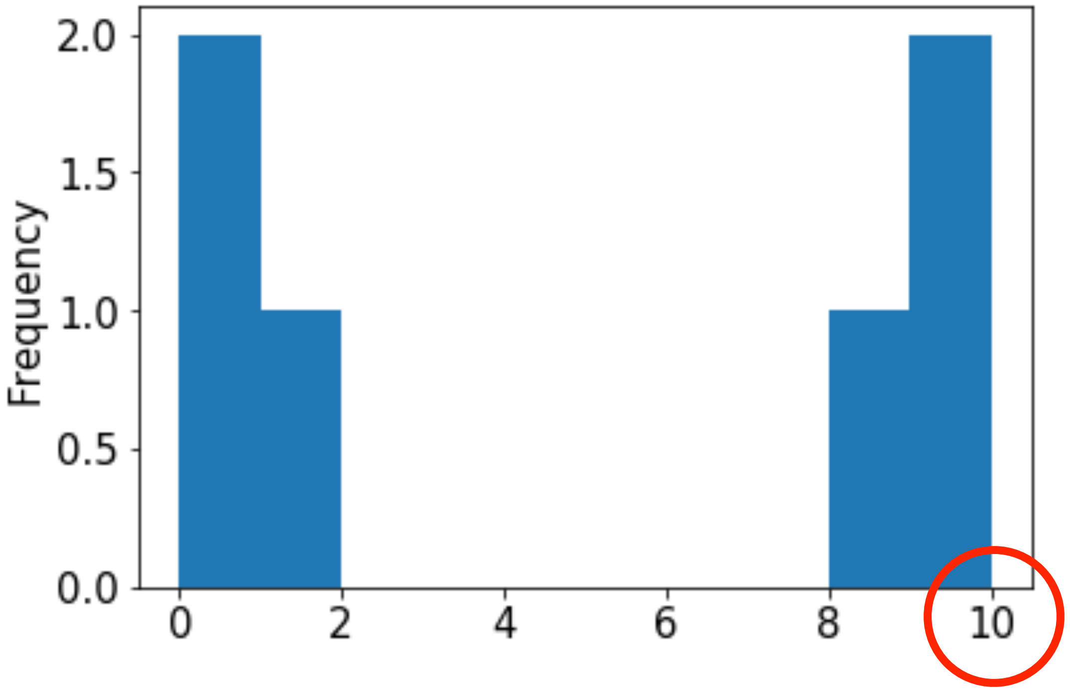

we can override the defaults

Frequencies across numbers

histograms are a good way to view frequencies across numbers

```
s = Series([0.1, 0, 1, 8, 9, 9.2])
```

~~`s.value_counts().sort_index().plot.bar()`~~

```
s.plot.hist(bins=range(11))
```

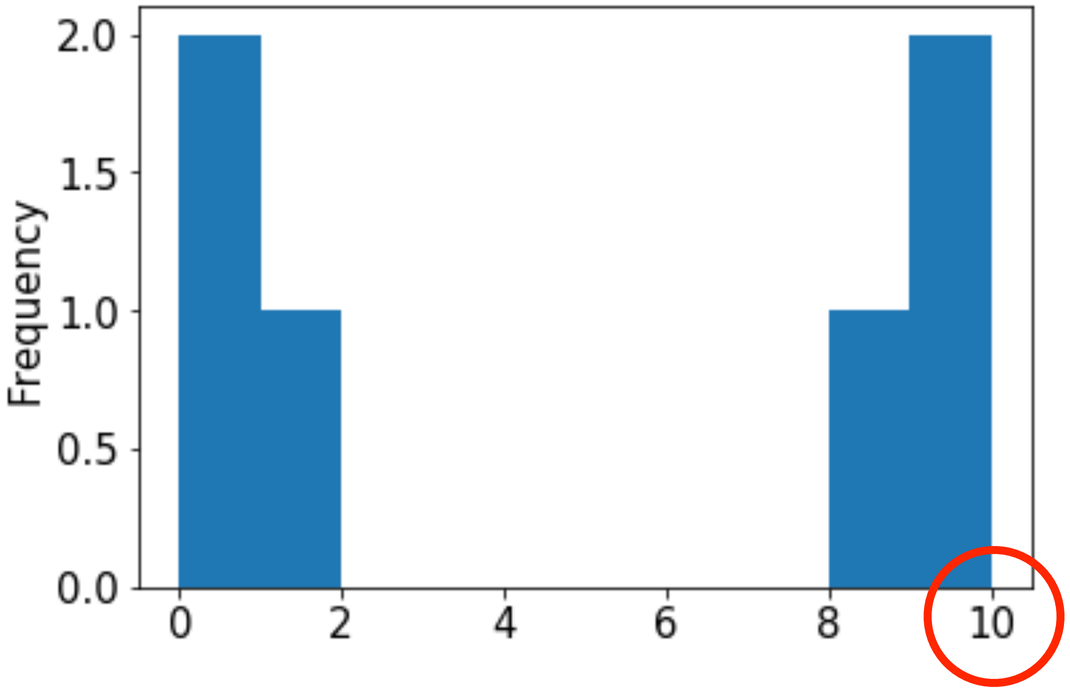

this is easily done with range

Demo: Visualize CoinSim Results

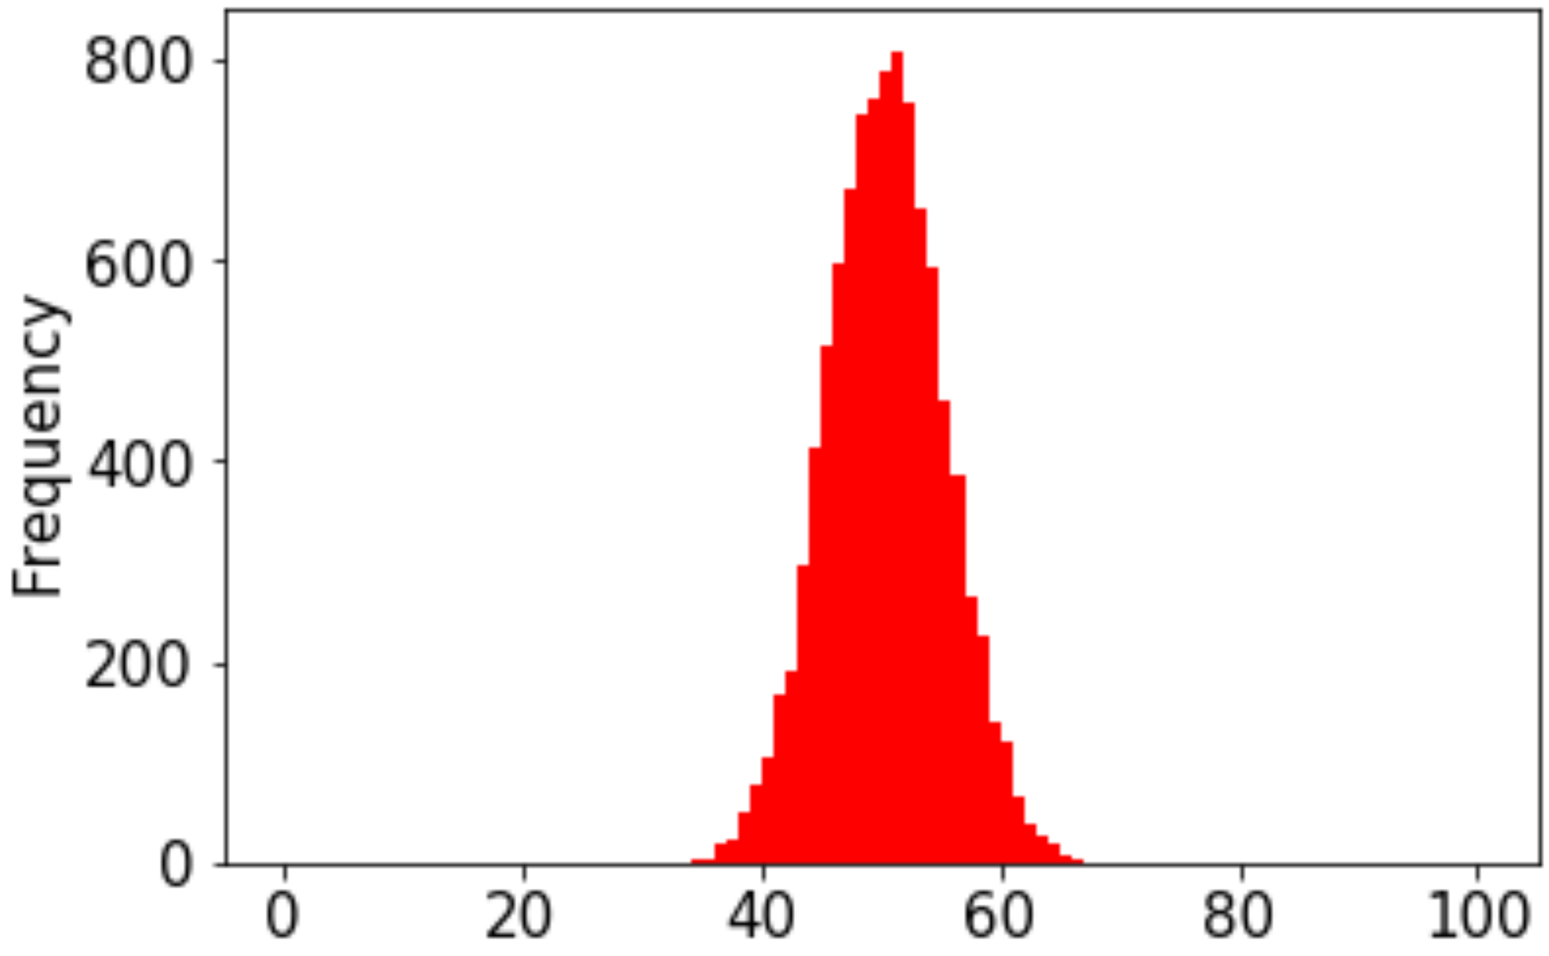

number of heads (out of 100)

Demo: Visualize CoinSim Results

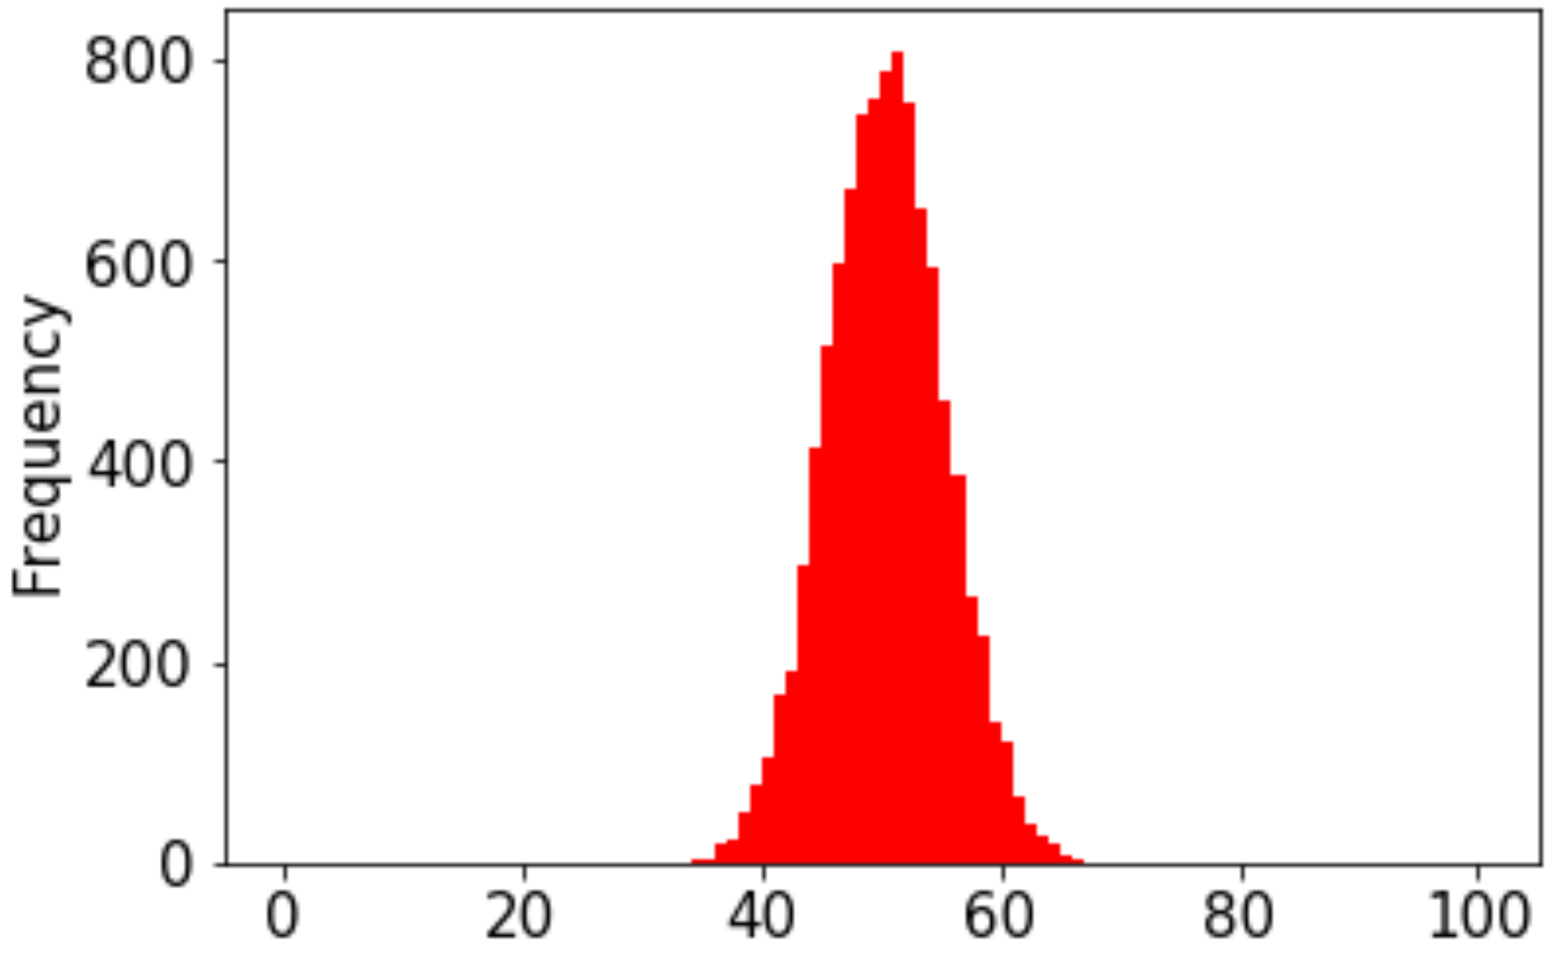

number of heads (out of 100)

this shape resembles what we often call a normal distribution or a "bell curve"

Demo: Visualize CoinSim Results

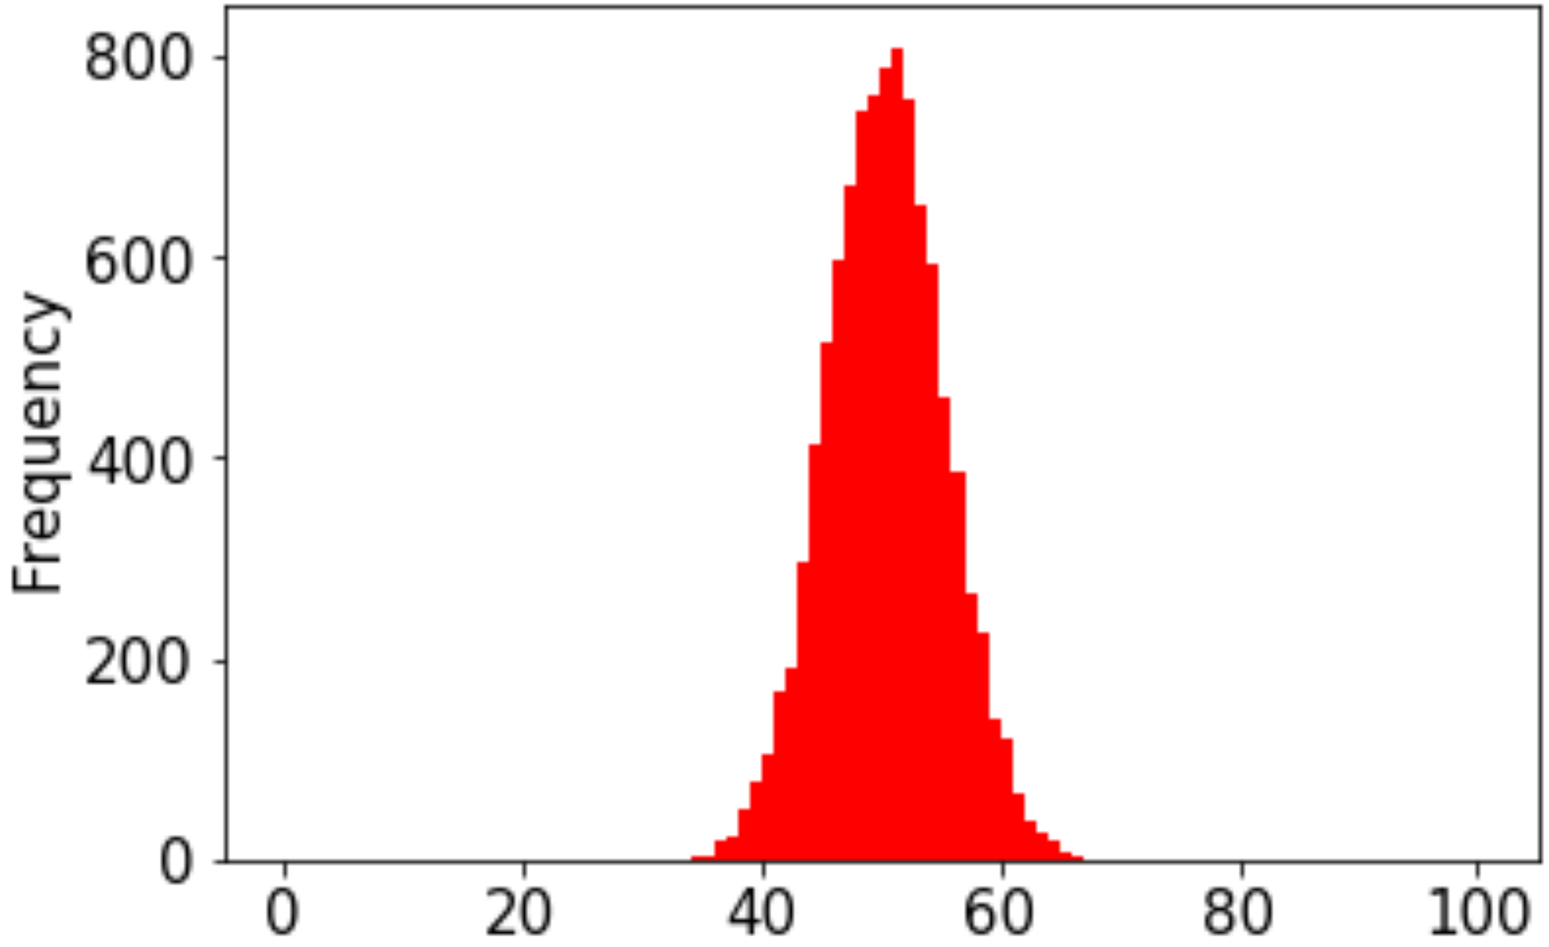

number of heads (out of 100)

this shape resembles what we often call a normal distribution or a "bell curve"

in general, if we take large samples enough times, the results will look like this
(we won't discuss exceptions here)

Demo: Visualize CoinSim Results

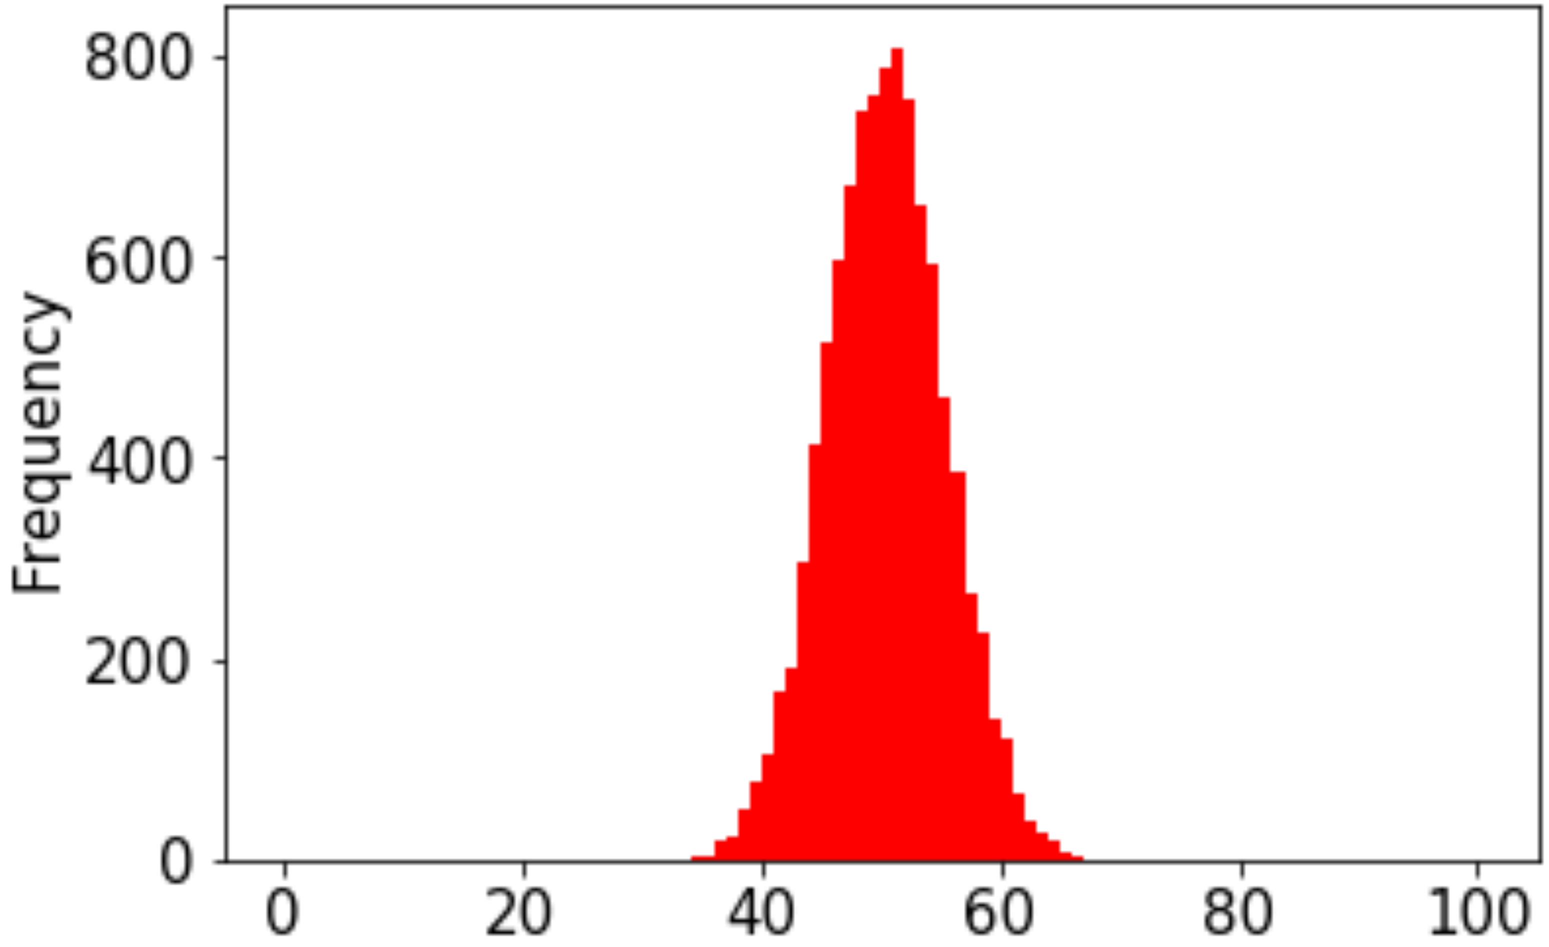

number of heads (out of 100)

numpy can directly generate random numbers fitting a normal distribution

this shape resembles what we often call a normal distribution or a "bell curve"

in general, if we take large samples enough times, the results will look like this
(we won't discuss exceptions here)

Outline

choice()

bugs and seeding

significance

histograms

normal()

normal

```
from numpy.random import choice, normal
import numpy as np

for i in range(10):
    print(normal())
```

normal

```
from numpy.random import choice, normal
import numpy as np

for i in range(10):
    print(normal())
```

average is 0 (over many calls)

numbers closer to 0 more likely

-x just as likely as x

Output:

```
-0.18638553993371157
0.02888452916769247
1.247456111372642З
-0.5388224399358179
-0.45143322136388525
-1.4001861112018241
0.2811937151186804７
0.260886189855659７
-0.19246288728955144
0.2979572961710292
```

normal

```
from numpy.random import choice, normal
import numpy as np

s = Series(normal(size=10000))
```

normal

```
from numpy.random import choice, normal
import numpy as np

s = Series(normal(size=10000))

s.plot.hist()
```

normal

```
from numpy.random import choice, normal
import numpy as np

s = Series(normal(size=10000))

s.plot.hist()
```

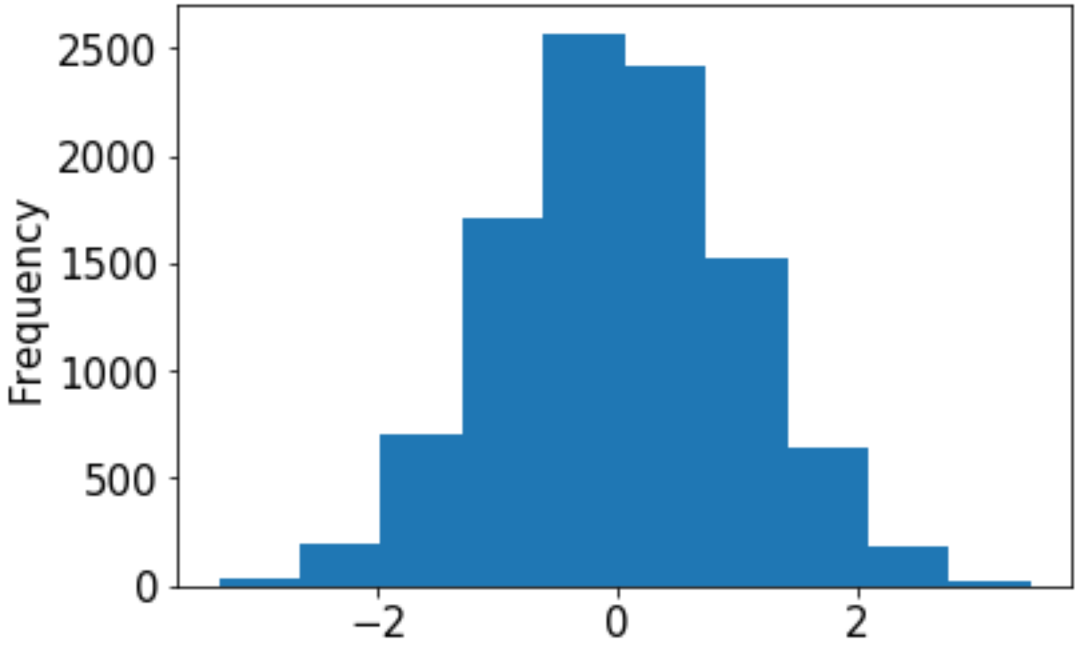

normal

```
from numpy.random import choice, normal
import numpy as np

s = Series(normal(size=10000))

s.plot.hist(bins=100)
```

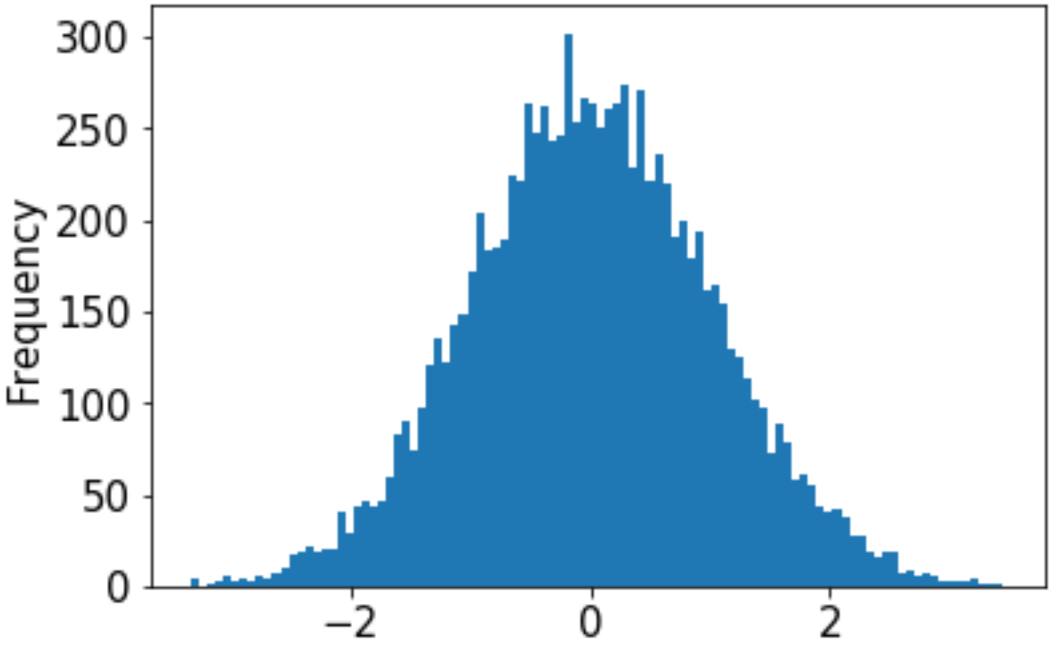

normal

```
from numpy.random import choice, normal
import numpy as np

s = Series(normal(size=10000))

s.plot.hist(bins=100, loc=[    ], scale=[    ])
```

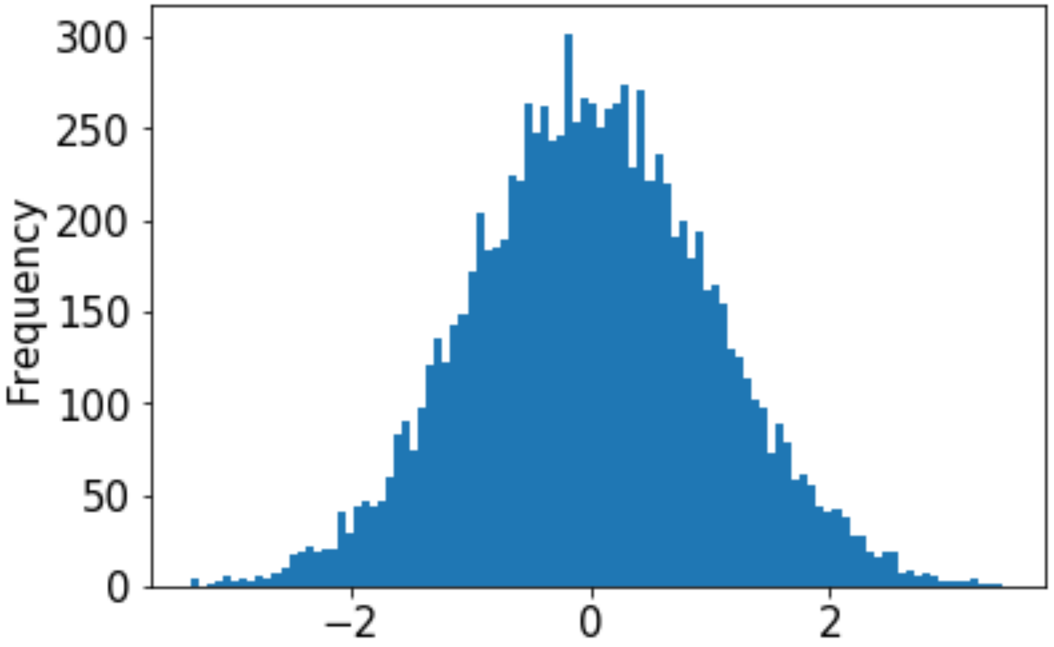

normal

```
from numpy.random import choice, normal
import numpy as np

s = Series(normal(size=10000))

s.plot.hist(bins=100, loc=[    ], scale=[    ])
```

try plugging in different values
(defaults are 0 and 1, respectively)

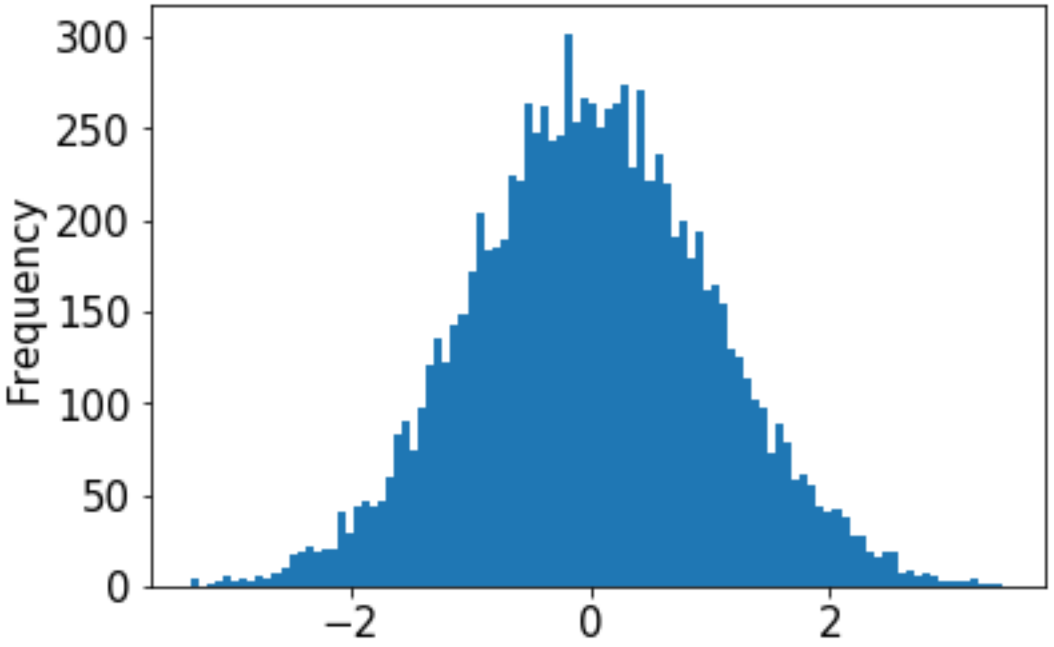

Demo: plot overlay

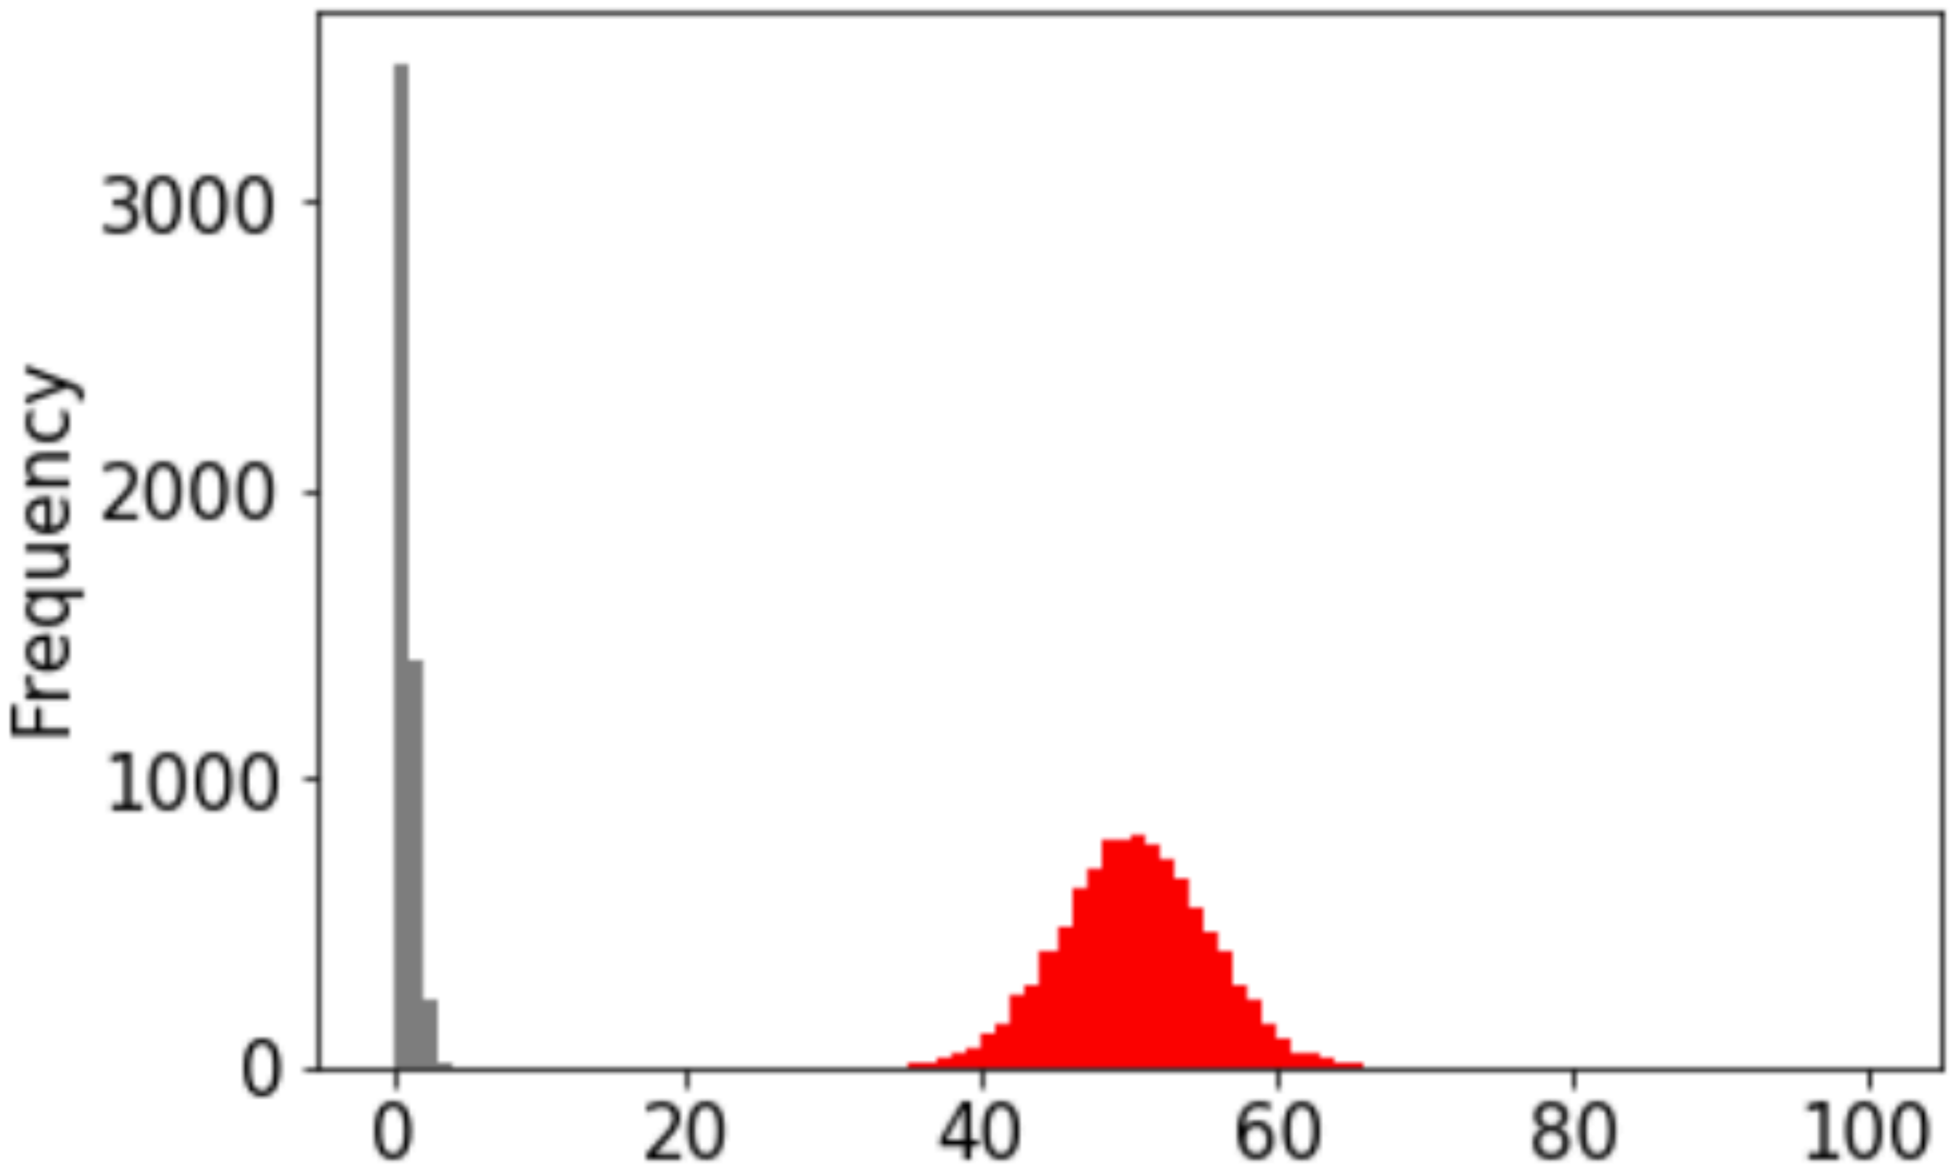

Demo: plot overlay

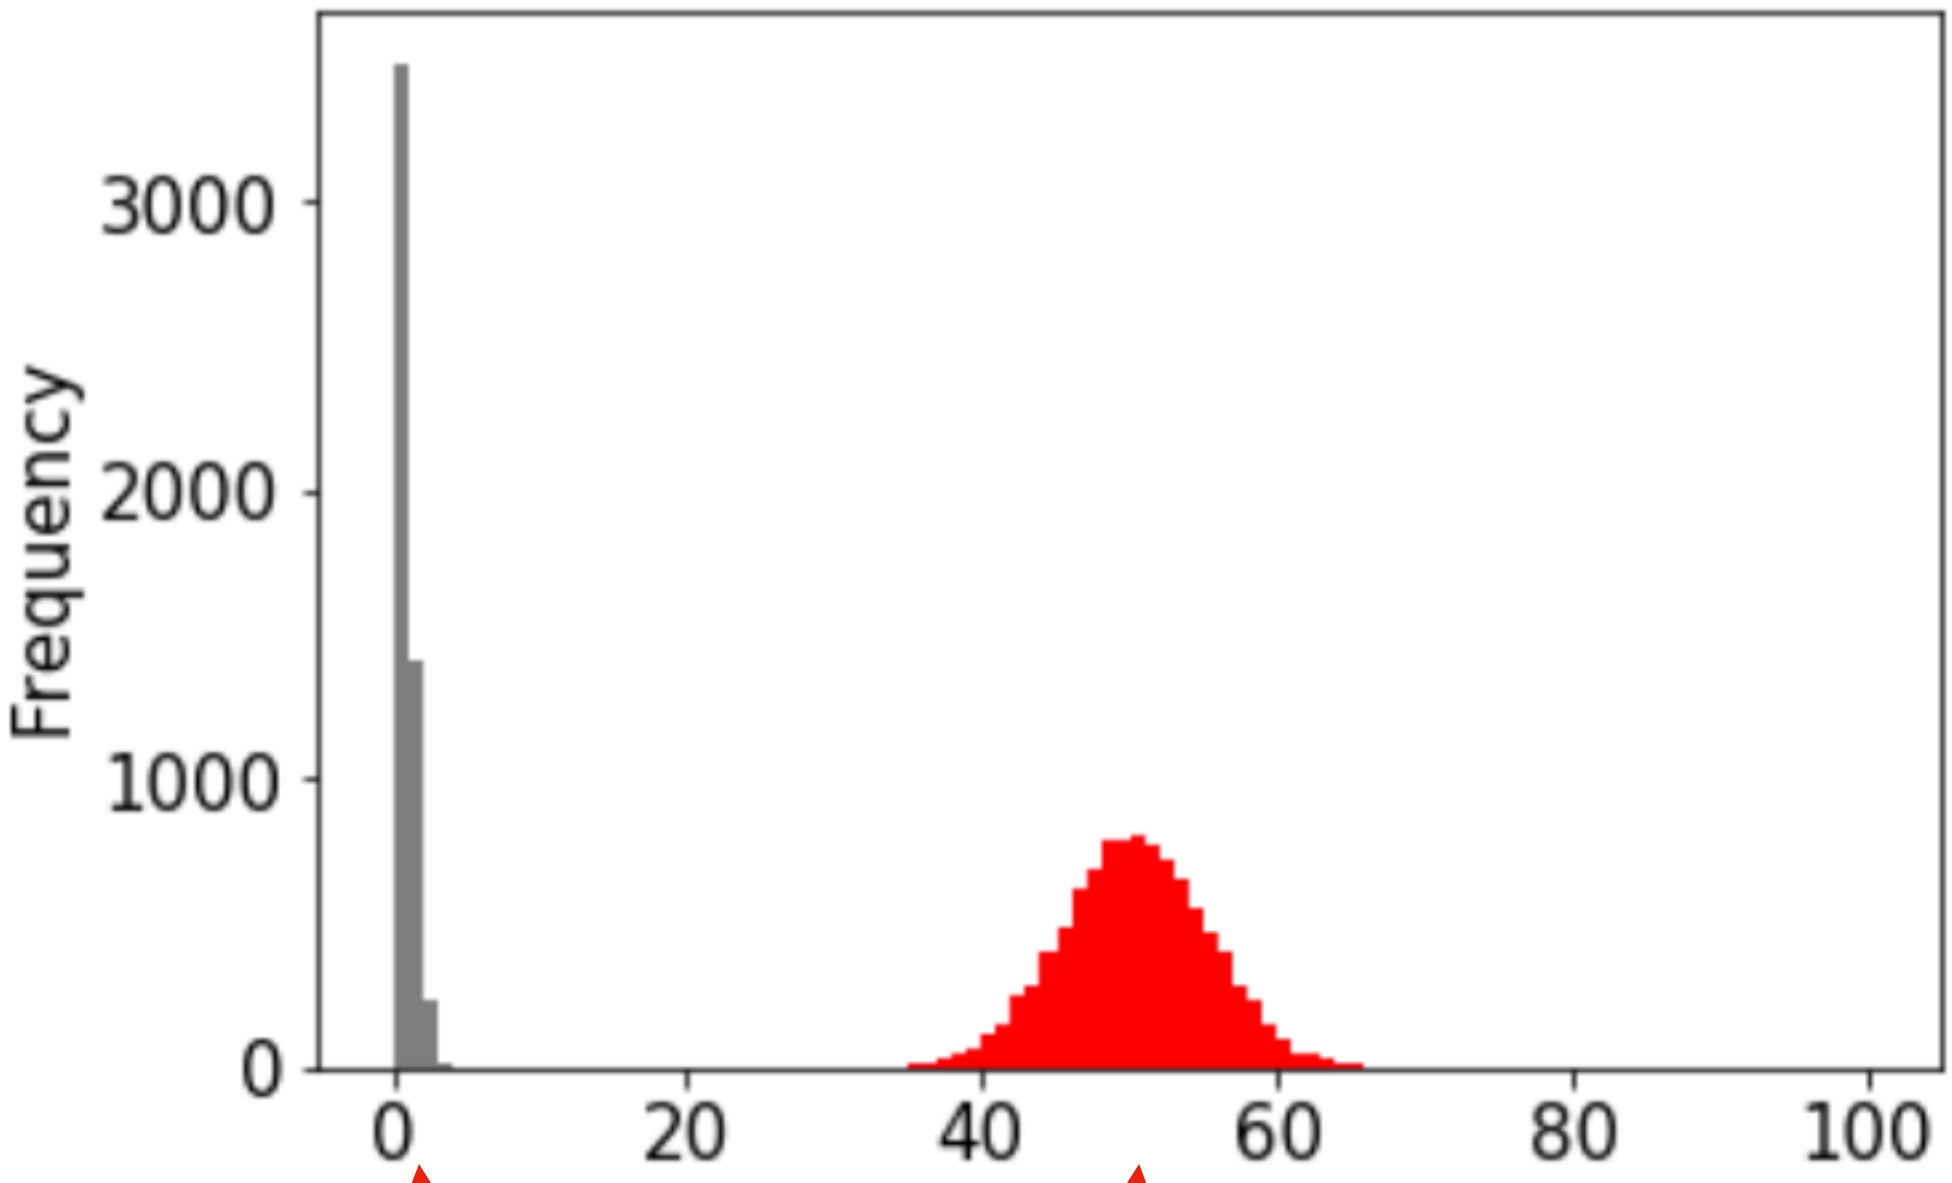

10K samples of 100 coin flips

10K samples from normal(size=10000)

Demo: plot overlay

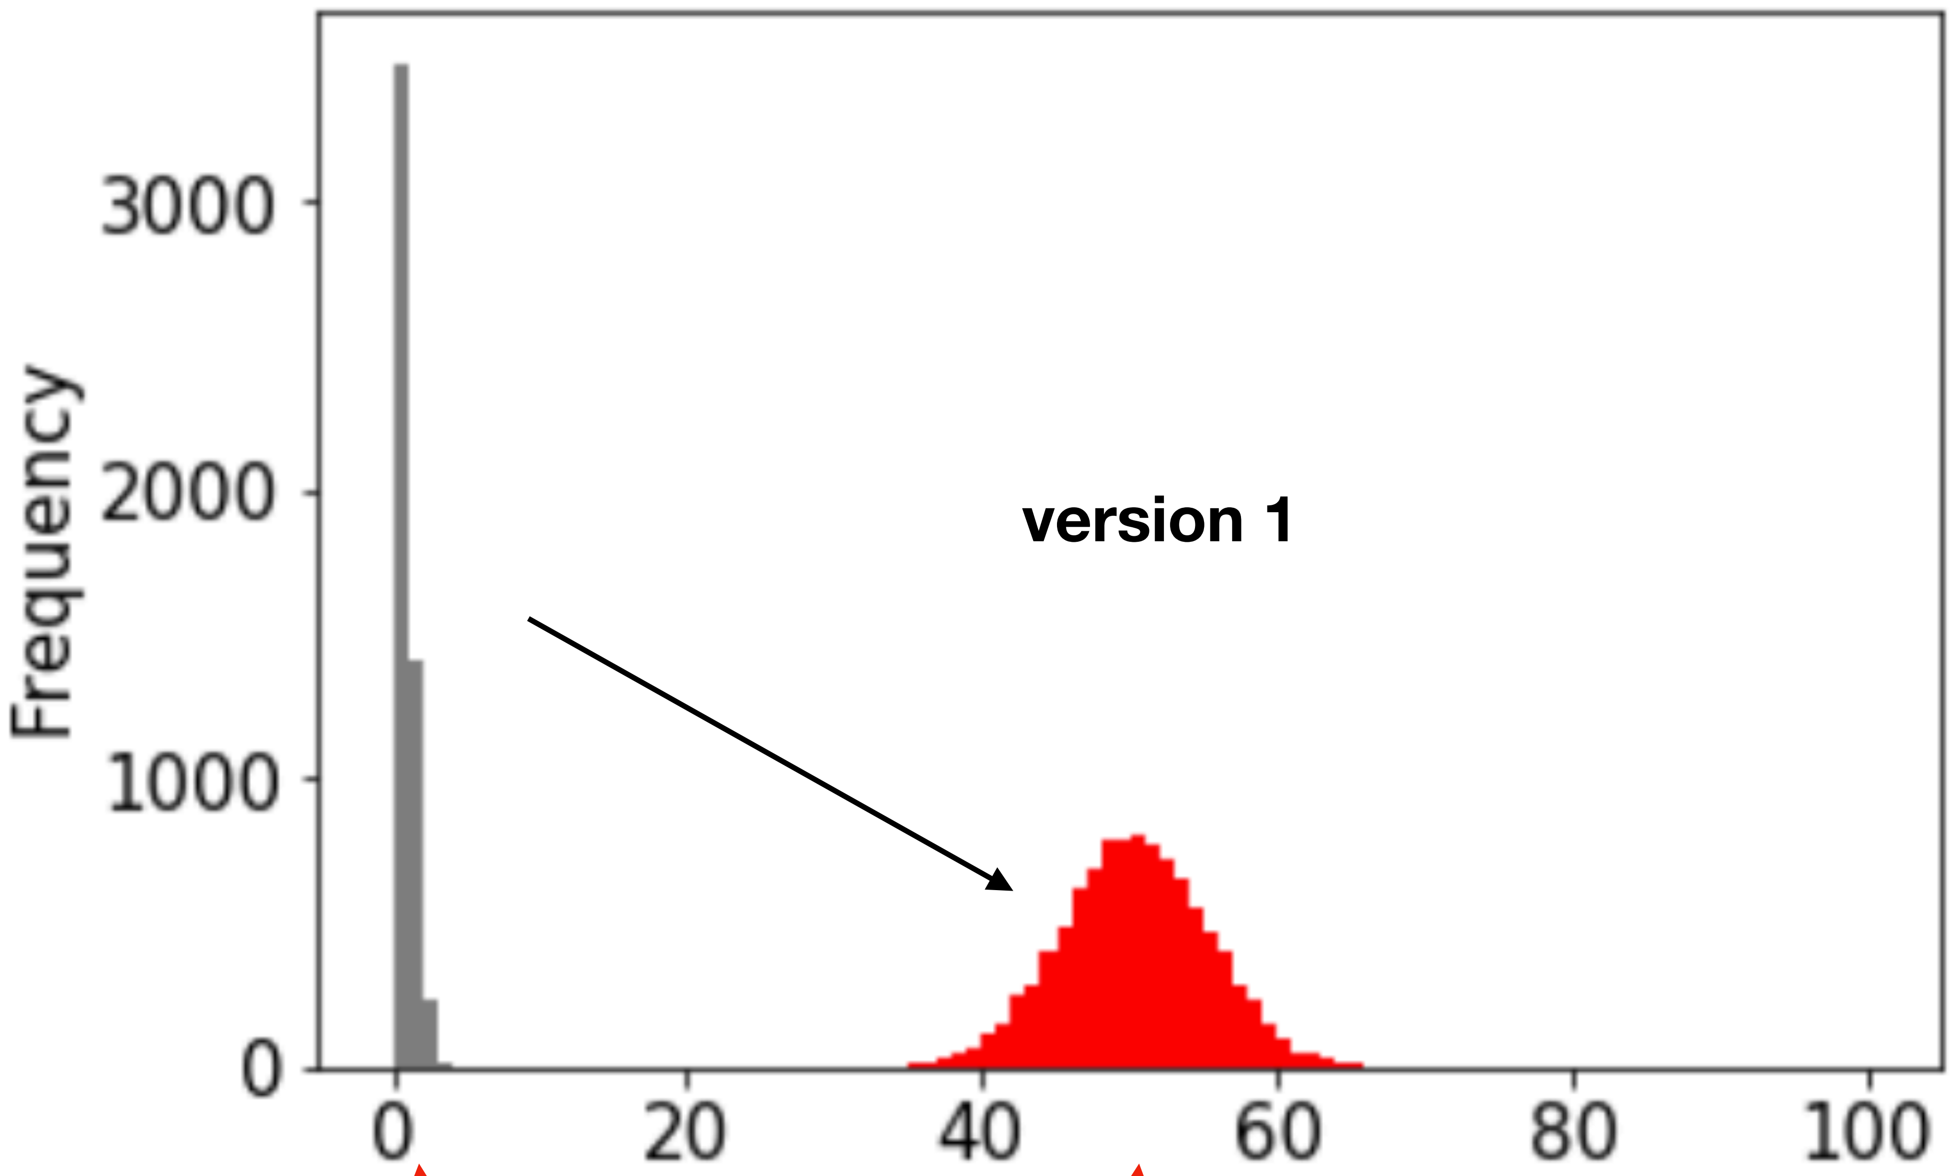

10K samples of 100 coin flips

10K samples from normal(size=10000)

goal: play with **`loc`** and **`scale`** arguments to normal until gray overlaps red

Demo: plot overlay

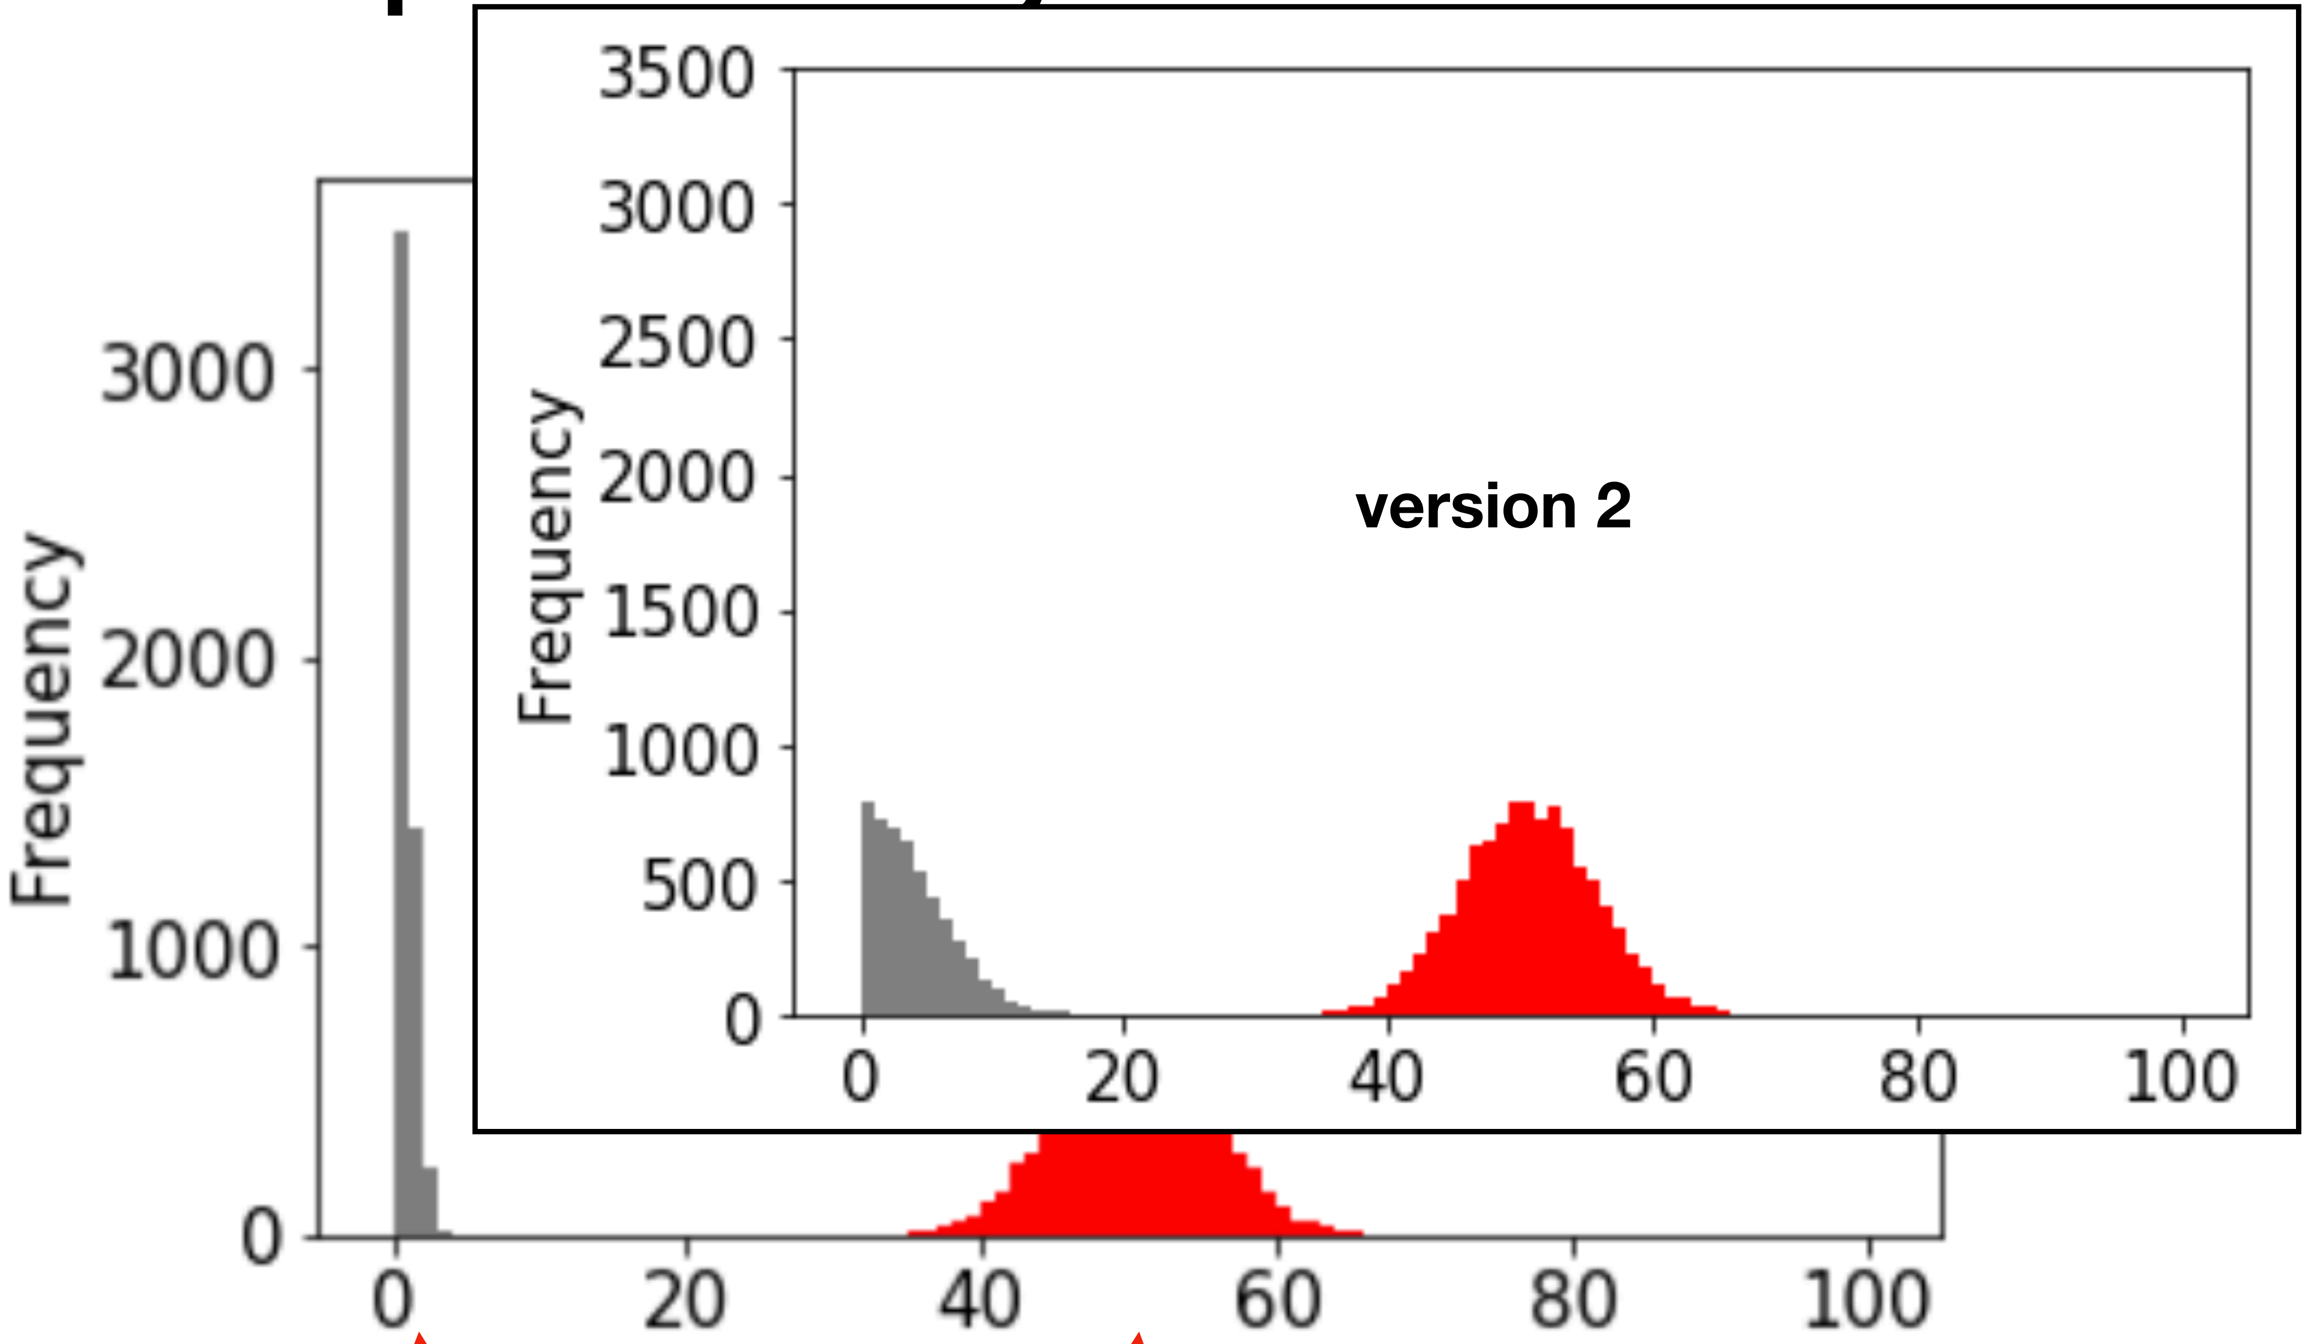

10K samples of 100 coin flips

10K samples from normal(size=10000)

goal: play with **`loc`** and **`scale`** arguments to normal until gray overlaps red

Demo: plot overlay

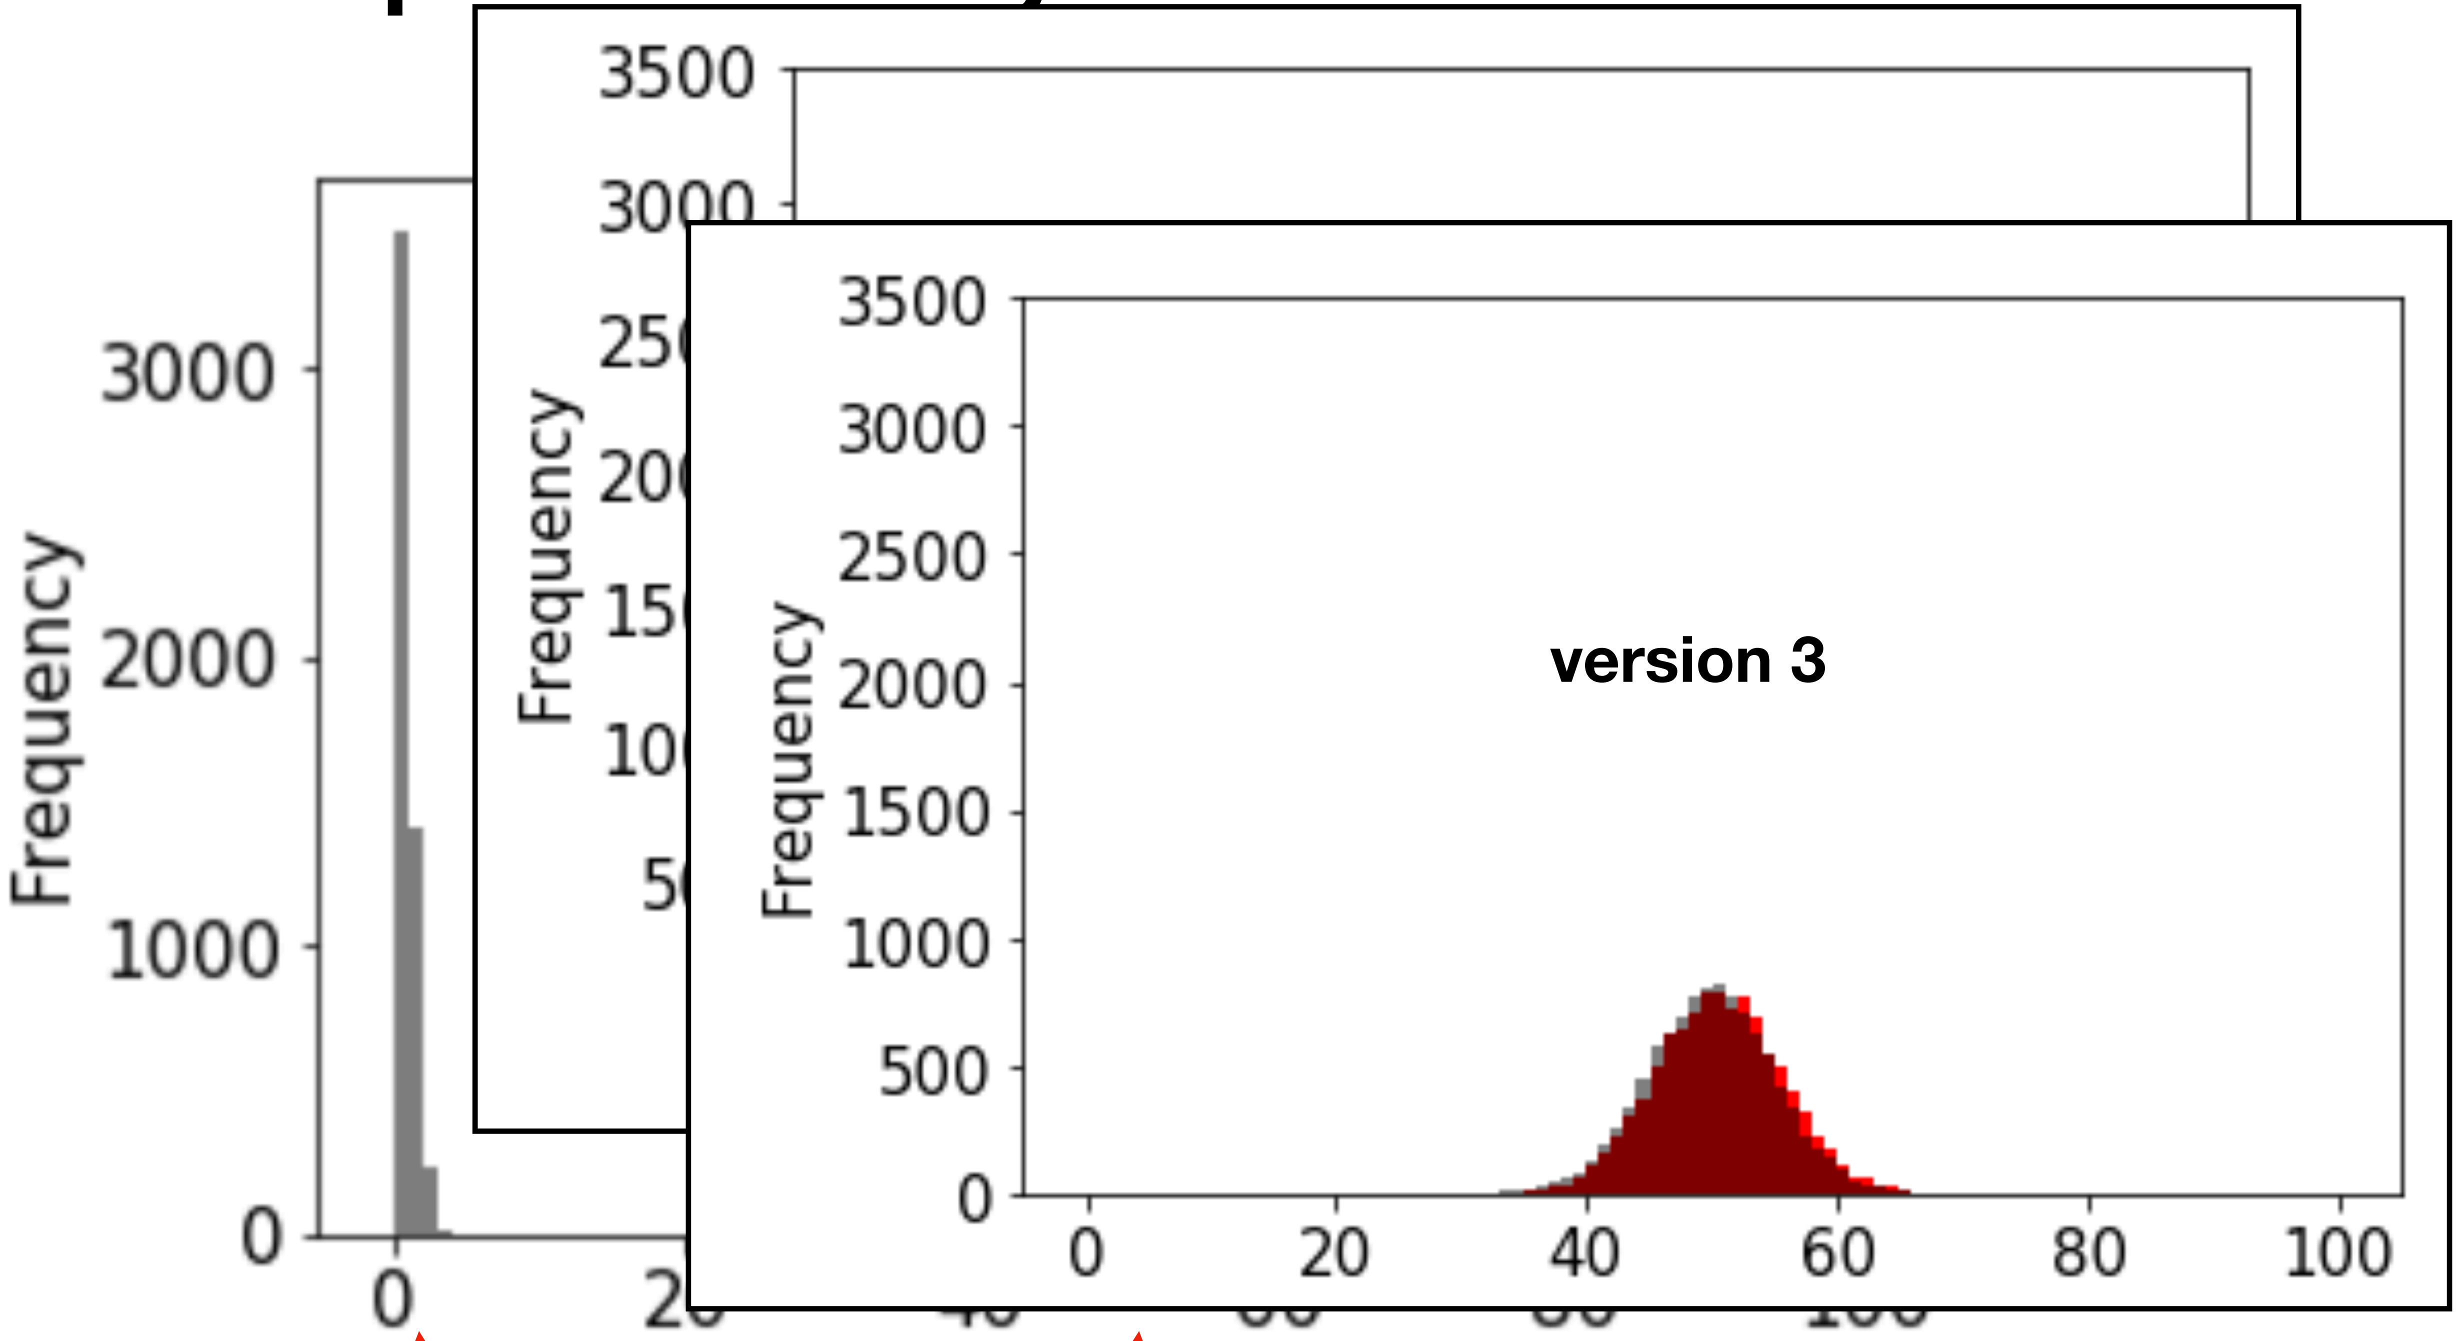

10K samples of 100 coin flips

10K samples from normal(size=10000)

goal: play with **`loc`** and **`scale`** arguments to normal until gray overlaps red